

CS 33

Data Representation (Part 3)

Byte Ordering

- **Four-byte integer**
 - 0x76543210
- **Stored at location 0x100**
 - which byte is at 0x100?
 - which byte is at 0x103?




| | | | |
|-------|-------|-------|-------|
| 10 | 32 | 54 | 76 |
| 0x100 | 0x101 | 0x102 | 0x103 |


Little-endian

| | | | |
|-------|-------|-------|-------|
| 76 | 54 | 32 | 10 |
| 0x100 | 0x101 | 0x102 | 0x103 |

Big-endian



Byte Ordering (2)



Quiz 1

```
int main() {  
    long x=1;  
    func((int *) &x);  
    return 0;  
}  
  
void func(int *arg) {  
    printf("%d\n", *arg);  
}
```

**What value is printed
on a big-endian 64-bit
computer?**

- a) 1
- b) 0
- c) 2^{32}
- d) $2^{32}-1$

Which Byte Ordering Do We Use?

```
int main() {  
    unsigned int x = 0x03020100;  
    unsigned char *xarray = (unsigned char *) &x;  
    for (int i=0; i<4; i++) {  
        printf("%02x", xarray[i]);  
    }  
    printf("\n");  
    return 0;  
}
```

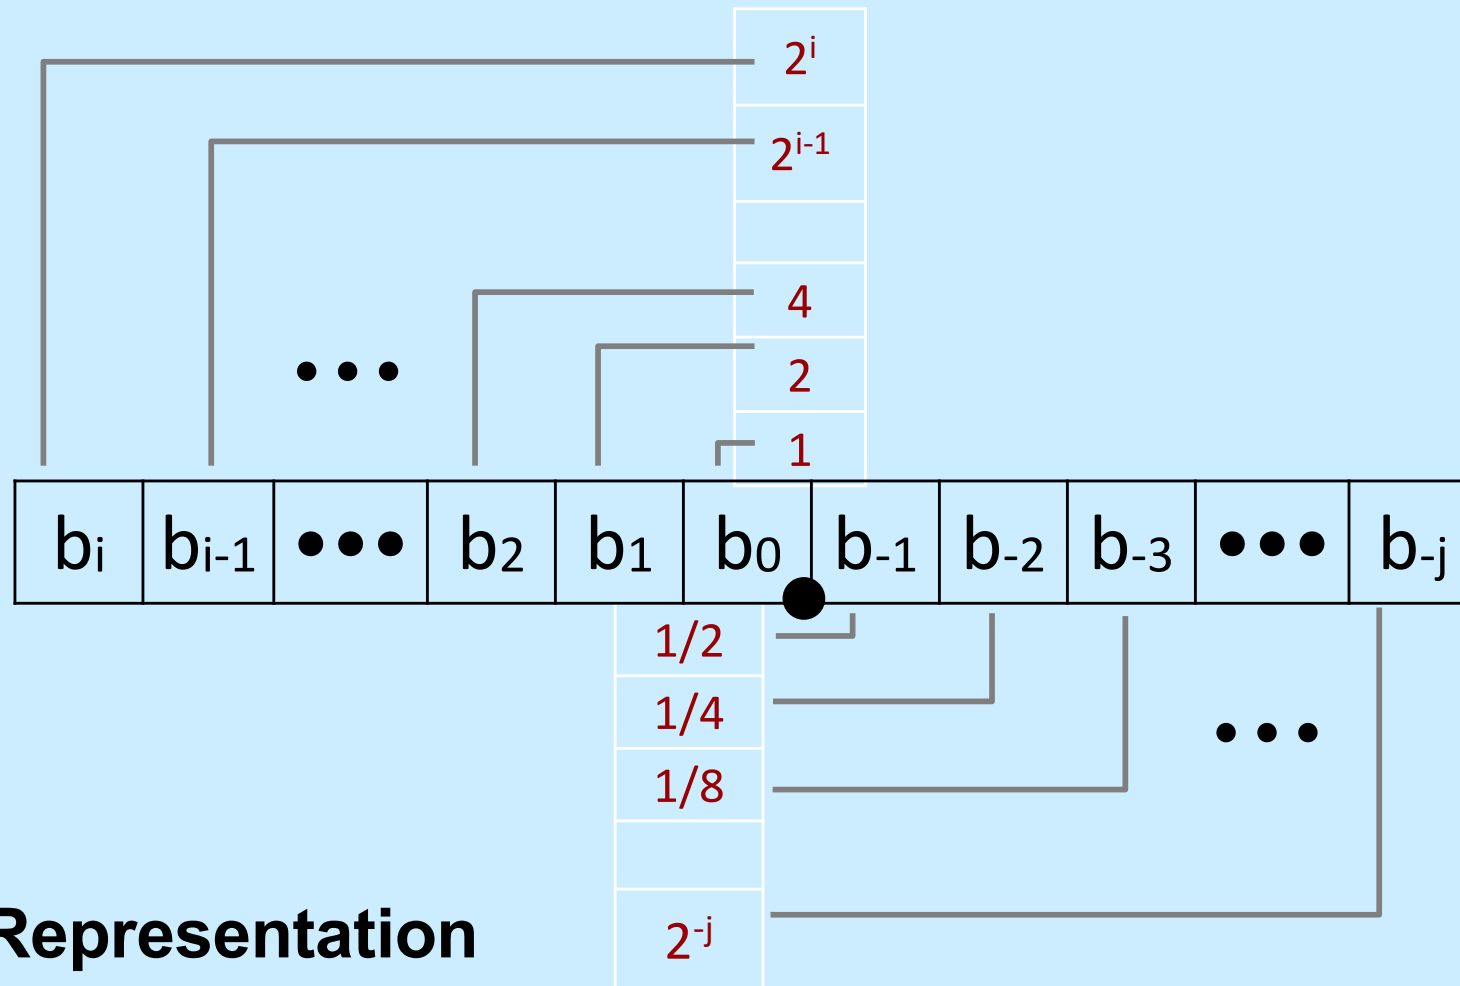
Possible results:

00010203
03020100

Fractional binary numbers

- What is 1011.101_2 ?

Fractional Binary Numbers



- **Representation**

- bits to right of “binary point” represent fractional powers of 2
- represents rational number:

$$\sum_{k=-j}^i b_k \times 2^k$$

Representable Numbers

- **Limitation #1**

- can exactly represent only numbers of the form $n/2^k$

- » other rational numbers have repeating bit representations

- value representation

- » 1/3 0.0101010101[01]...₂

- » 1/5 0.001100110011[0011]...₂

- » 1/10 0.0001100110011[0011]...₂

- **Limitation #2**

- just one setting of decimal point within the w bits

- » limited range of numbers (very small values? very large?)

IEEE Floating Point

- **IEEE Standard 754**
 - established in 1985 as uniform standard for floating point arithmetic
 - » before that, many idiosyncratic formats
 - supported on all major CPUs
- **Driven by numerical concerns**
 - nice standards for rounding, overflow, underflow
 - hard to make fast in hardware
 - » numerical analysts predominated over hardware designers in defining standard

Floating-Point Representation

- Numerical Form:

$$(-1)^s M 2^E$$

- sign bit **s** determines whether number is negative or positive
- significand **M** normally a fractional value in range $[1.0, 2.0)$
- exponent **E** weights value by power of two

- Encoding

- MSB **s** is sign bit **s**
- exp field encodes **E** (but is not equal to E)
- frac field encodes **M** (but is not equal to M)



Precision options

- **Single precision: 32 bits**



- **Double precision: 64 bits**



- **Extended precision: 80 bits (Intel only)**



“Normalized” Values

- When: $\text{exp} \neq 000\dots 0$ and $\text{exp} \neq 111\dots 1$
- Exponent coded as biased value: $E = \text{Exp} - \text{Bias}$
 - exp : unsigned value exp
 - $\text{bias} = 2^{k-1} - 1$, where k is number of exponent bits
 - » single precision: 127 (Exp: 1...254, E: -126...127)
 - » double precision: 1023 (Exp: 1...2046, E: -1022...1023)
- Significand coded with implied leading 1: $M = 1.\text{xxx}\dots\text{x}_2$
 - $\text{xxx}\dots\text{x}$: bits of frac
 - minimum when $\text{frac} = 000\dots 0$ ($M = 1.0$)
 - maximum when $\text{frac} = 111\dots 1$ ($M = 2.0 - \epsilon$)
 - get extra leading bit for “free”

Normalized Encoding Example

- **Value:** `float F = 15213.0;`

$$\begin{aligned} - 15213_{10} &= 11101101101101_2 \\ &= 1.1101101101101_2 \times 2^{13} \end{aligned}$$


- **Significand**

$$\begin{aligned} M &= 1.\underline{1101101101101}_2 \\ \text{frac} &= \underline{1101101101101}0000000000_2 \end{aligned}$$

- **Exponent**

$$\begin{aligned} E &= 13 \\ \text{bias} &= 127 \\ \text{exp} &= 140 = 10001100_2 \end{aligned}$$

- **Result:**




Denormalized Values

- Condition: $\text{exp} = 000\dots 0$
- Exponent value: $E = -\text{Bias} + 1$ (instead of $E = 0 - \text{Bias}$)
- Significand coded with implied leading 0:
 $M = 0.\text{xxx}\dots\text{x}_2$
 - $\text{xxx}\dots\text{x}$: bits of frac
- Cases
 - $\text{exp} = 000\dots 0, \text{frac} = 000\dots 0$
 - » represents zero value
 - » note distinct values: $+0$ and -0 (why?)
 - $\text{exp} = 000\dots 0, \text{frac} \neq 000\dots 0$
 - » numbers closest to 0.0
 - » equispaced


Special Values

- **Condition: $\text{exp} = 111\dots 1$**
 - **Case: $\text{exp} = 111\dots 1, \text{frac} = 000\dots 0$**
 - represents value ∞ (infinity)
 - operation that overflows
 - both positive and negative
 - e.g., $1.0/0.0 = -1.0/-0.0 = +\infty$, $1.0/-0.0 = -\infty$
 - **Case: $\text{exp} = 111\dots 1, \text{frac} \neq 000\dots 0$**
 - not-a-number (NaN)
 - represents case when no numeric value can be determined
 - e.g., $\text{sqrt}(-1)$, $\infty - \infty$, $\infty \times 0$
-

Visualization: Floating-Point Encodings



Tiny Floating-Point Example



- **8-bit Floating Point Representation**
 - the sign bit is in the most significant bit
 - the next four bits are the exponent, with a bias of 7
 - the last three bits are the *frac*
- **Same general form as IEEE Format**
 - normalized, denormalized
 - representation of 0, NaN, infinity

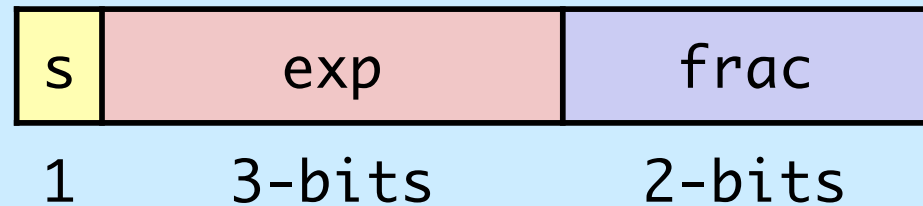
Dynamic Range (Positive Only)

| | s | exp | frac | E | Value | |
|----------------------|-----|------|------|-----|----------------------|--------------------|
| Denormalized numbers | 0 | 0000 | 000 | -6 | 0 | |
| | 0 | 0000 | 001 | -6 | $1/8 * 1/64 = 1/512$ | closest to zero |
| | 0 | 0000 | 010 | -6 | $2/8 * 1/64 = 2/512$ | |
| | ... | | | | | |
| | 0 | 0000 | 110 | -6 | $6/8 * 1/64 = 6/512$ | |
| | 0 | 0000 | 111 | -6 | $7/8 * 1/64 = 7/512$ | largest denorm |
| Normalized numbers | 0 | 0001 | 000 | -6 | $8/8 * 1/64 = 8/512$ | smallest norm |
| | 0 | 0001 | 001 | -6 | $9/8 * 1/64 = 9/512$ | |
| | ... | | | | | |
| | 0 | 0110 | 110 | -1 | $14/8 * 1/2 = 14/16$ | |
| | 0 | 0110 | 111 | -1 | $15/8 * 1/2 = 15/16$ | closest to 1 below |
| | 0 | 0111 | 000 | 0 | $8/8 * 1 = 1$ | |
| | 0 | 0111 | 001 | 0 | $9/8 * 1 = 9/8$ | closest to 1 above |
| | 0 | 0111 | 010 | 0 | $10/8 * 1 = 10/8$ | |
| | ... | | | | | |
| | 0 | 1110 | 110 | 7 | $14/8 * 128 = 224$ | |
| | 0 | 1110 | 111 | 7 | $15/8 * 128 = 240$ | largest norm |
| | 0 | 1111 | 000 | n/a | inf | |


Distribution of Values

- 6-bit IEEE-like format

- e = 3 exponent bits
- f = 2 fraction bits
- bias is $2^{3-1}-1 = 3$





- Notice how the distribution gets denser toward zero.



Distribution of Values (close-up view)

- 6-bit IEEE-like format


- e = 3 exponent bits
- f = 2 fraction bits
- bias is 3



Quiz 2

- **6-bit IEEE-like format**

- e = 3 exponent bits
- f = 2 fraction bits
- bias is 3




What number is represented by 0 010 10?

- a) 3
- b) 1.5
- c) .75
- d) none of the above

Mapping Real Numbers to Float

- The real number 3 is represented as
0 100 10
- The real number 3.5 is represented as
0 100 11
- How is the real number 3.4 represented?
0 100 11
- How is the real number π represented?
0 100 10




Mapping Real Numbers to Float

- If R is a real number, it's mapped to the floating-point number whose value is closest to R
- What if it's midway between two values?
 - rounding rules determine outcome

Floats are Sets of Values

- If A, B, and C are successive floating-point values
 - e.g., 010001, 010010, and 010011
- B represents all real numbers from midway between A and B through midway between B and C



Significance

- **Normalized numbers**
 - for a particular exponent value E and an S -bit significand, the range from 2^E up to 2^{E+1} is divided into 2^S equi-spaced floating-point values
 - » thus each floating-point value represents $1/2^S$ of the range of values with that exponent
 - » all bits of the significand are important
 - » we say that there are S significant bits – for reasonably large S , each floating-point value covers a rather small part of the range
 - high accuracy
 - for $S=23$ (32-bit float), accurate to one in 2^{23} (.0000119% accuracy)

Significance

- **Unnormalized numbers**
 - high-order zero bits of the significand aren't important
 - in 8-bit floating point, 0 0000 001 represents 2^{-9}
 - » it is the only value with that exponent: 1 significant bit (either 2^{-9} or 0)
 - 0 0000 010 represents 2^{-8}
0 0000 011 represents $1.5 \cdot 2^{-8}$
 - » only two values with exponent -8: 2 significant bits (encoding those two values, as well as 2^{-9} and 0)
 - fewer significant bits mean less accuracy
 - 0 0000 001 represents a range of values from $.5 \cdot 2^{-9}$ to $1.5 \cdot 2^{-9}$
 - 50% accuracy


+/- Zero

- **Only one zero for ints**
 - an int is a single number, not a range of numbers, thus there can be only zero
- **Floating-point zero**
 - a range of numbers around the real 0
 - it really matters which side of 0 we're on!
 - » a very large negative number divided by a very small negative number should be positive
$$-\infty / -0 = +\infty$$
 - » a very large positive number divided by a very small negative number should be negative
$$+\infty / -0 = -\infty$$

CS 33

Intro to Machine Programming

Machine Model



Memory



Instructions

Data

The diagram on the left consists of two stacked, light-orange rounded rectangles. The top rectangle is labeled 'Instructions' and the bottom rectangle is labeled 'Data'. This represents a memory architecture where instructions and data are stored in separate locations.


or



**Instructions
are Data**

The diagram on the right is a single, tall, light-orange rounded rectangle. Inside, the text 'Instructions are Data' is centered. This represents a memory architecture where instructions and data are treated as a single, unified entity.

Processor: Some Details



Processor: Basic Operation

```
while (forever) {  
  fetch instruction IP points at  
  decode instruction  
  fetch operands  
  execute  
  store results  
  update IP and condition code  
}
```


Instructions ...

| | | | |
|----------------|-----------------|-----------------|------------|
| Op code | Operand1 | Operand2 | ... |
|----------------|-----------------|-----------------|------------|

Operands

- **Form**
 - immediate vs. reference
 - » value vs. address
- **How many?**
 - 3
 - » add a,b,c
 - $c = a + b$
 - 2
 - » add a,b
 - $b += a$

Operands (continued)

- **Accumulator**
 - special memory in the processor
 - » known as a *register*
 - » fast access
 - allows single-operand instructions
 - » add a
 - `acc += a`
 - » add b
 - `acc += b`

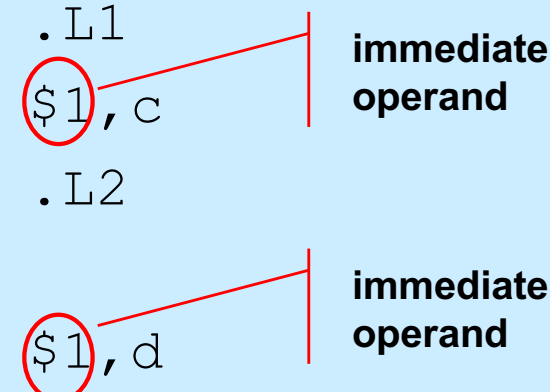
From C to Assembler ...

```
a = (b + c) * d;
```

```
mov    b,%acc  
add    c,%acc  
mul    d,%acc  
mov    %acc,a
```

```
if (a<b)  
    c = 1;  
else  
    d = 1;
```

```
cmp    a,b  
jge    .L1  
mov    $1,c  
jmp    .L2  
.L1  
mov    $1,d  
.L2
```



immediate operand

immediate operand

Condition Codes

- **Set of flags giving status of most recent operation:**
 - **zero flag**
 - » result was zero
 - **sign flag**
 - » for signed arithmetic interpretation: sign bit is set
 - **overflow flag**
 - » for signed arithmetic interpretation
 - **carry flag (generated by carry or borrow out of most-significant bit)**
 - » for unsigned arithmetic interpretation
- **Set implicitly by arithmetic instructions**
- **Set explicitly by compare instruction**
 - **cmp a,b**
 - » sets flags based on result of $b-a$

Examples (1)

- **Assume 32-bit arithmetic**
- **x is 0x80000000**
 - **TMIN** if interpreted as two's-complement
 - **2^{31}** if interpreted as unsigned
- **x-1 (0x7fffffff)**
 - **TMAX** if interpreted as two's-complement
 - **$2^{31}-1$** if interpreted as unsigned
 - **zero flag** is not set
 - **sign flag** is not set
 - **overflow flag** is set
 - **carry flag** is not set

Examples (2)

- **x is 0xffffffff**
 - -1 if interpreted as two's-complement
 - UMAX ($2^{32}-1$) if interpreted as unsigned
- **x+1 (0x00000000)**
 - zero under either interpretation
 - zero flag is set
 - sign flag is not set
 - overflow flag is not set
 - carry flag is set

Examples (3)

- **x is 0xffffffff**
 - -1 if interpreted as two's-complement
 - UMAX ($2^{32}-1$) if interpreted as unsigned
- **x+2 (0x00000001)**
 - (+)1 under either interpretation
 - zero flag is not set
 - sign flag is not set
 - overflow flag is not set
 - carry flag is set

Quiz 3

- **Set of flags giving status of most recent operation:**
 - zero flag
 - » result was zero
 - sign flag
 - » for signed arithmetic interpretation: sign bit is set
 - overflow flag
 - » for signed arithmetic interpretation
 - carry flag (generated by carry or borrow out of most-significant bit)
 - » for unsigned arithmetic interpretation
- **Set explicitly by compare instruction**
 - `cmp a,b`
 - » sets flags based on result of `b-a`

Which flags are set to one by “`cmp 2,1`”?

- a) overflow flag only**
- b) carry flag only**
- c) sign and carry flags only**
- d) sign and overflow flags only**
- e) sign, overflow, and carry flags**

Jump Instructions

- **Unconditional jump**
 - just do it
- **Conditional jump**
 - to jump or not to jump determined by condition-code flags
 - field in the op code indicates how this is computed
 - in assembler language, simply say
 - » **je**
 - jump on equal
 - » **jne**
 - jump on not equal
 - » **jg**
 - jump on greater than (signed)
 - » **etc.**

Addresses


```
int a, b, c, d;
```

```
int main() {  
    a = (b + c) * d;  
    ...  
}
```

```
mov    b, %acc  
add    c, %acc  
mul    d, %acc  
mov    %acc, a
```

```
mov    1004, %acc  
add    1008, %acc  
mul    1012, %acc  
mov    %acc, 1000
```

1012:
1008:
1004:
1000:



Memory

Addresses

```
int b;
```

```
int func(int c, int d) {  
    int a;  
    a = (b + c) * d;  
    ...  
}
```


```
mov    ?, %acc  
add    ?, %acc  
mul    ?, %acc  
mov    %acc, ?
```

- One copy of *b* for duration of program's execution
 - *b*'s address is the same for each call to *func*
- Different copies of *a*, *c*, and *d* for each call to *func*
 - addresses are different in each call

Relative Addresses


- **Absolute address**
 - actual location in memory
- **Relative address**
 - offset from some other location

- Blob's absolute address is 10000
- Datum's relative address (to Blob) is 100
 - its absolute address is 10100



Base Registers

```
mov $10000, %base  
mov $10, 100(%base)
```

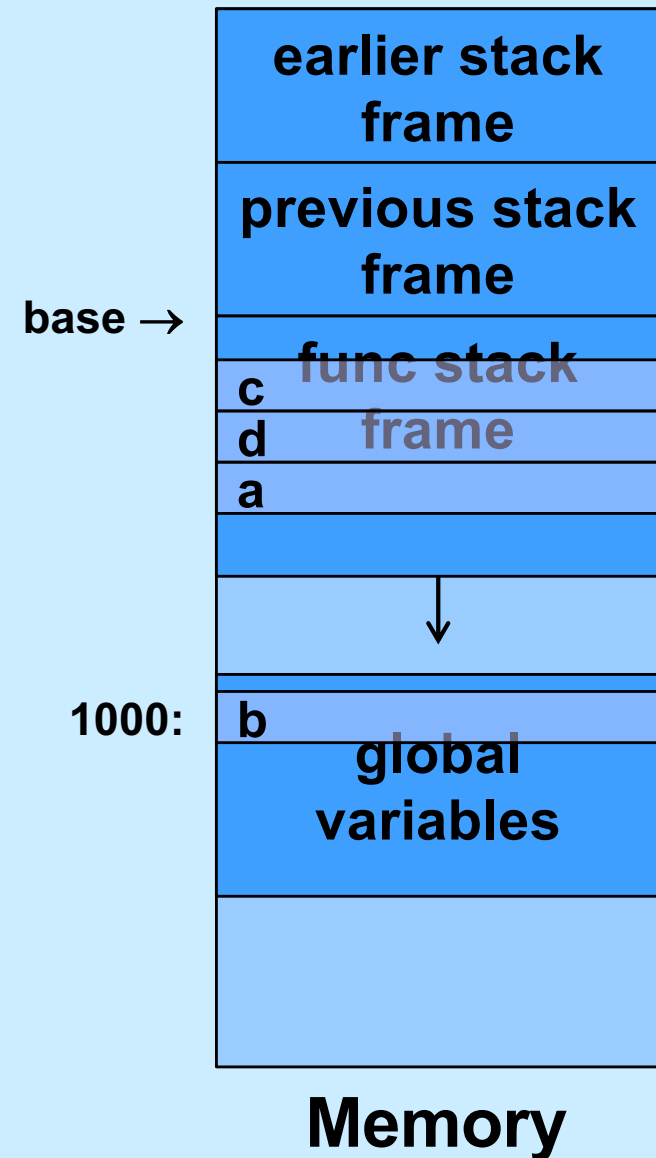


Addresses

```
int b;
```

```
int func(long c, long d) {  
    long a;  
    a = (b + c) * d;  
    ...  
}
```

```
mov    1000,%acc  
add    -8(%base),%acc  
mul    -16(%base),%acc  
mov    %acc,-24(%base)
```

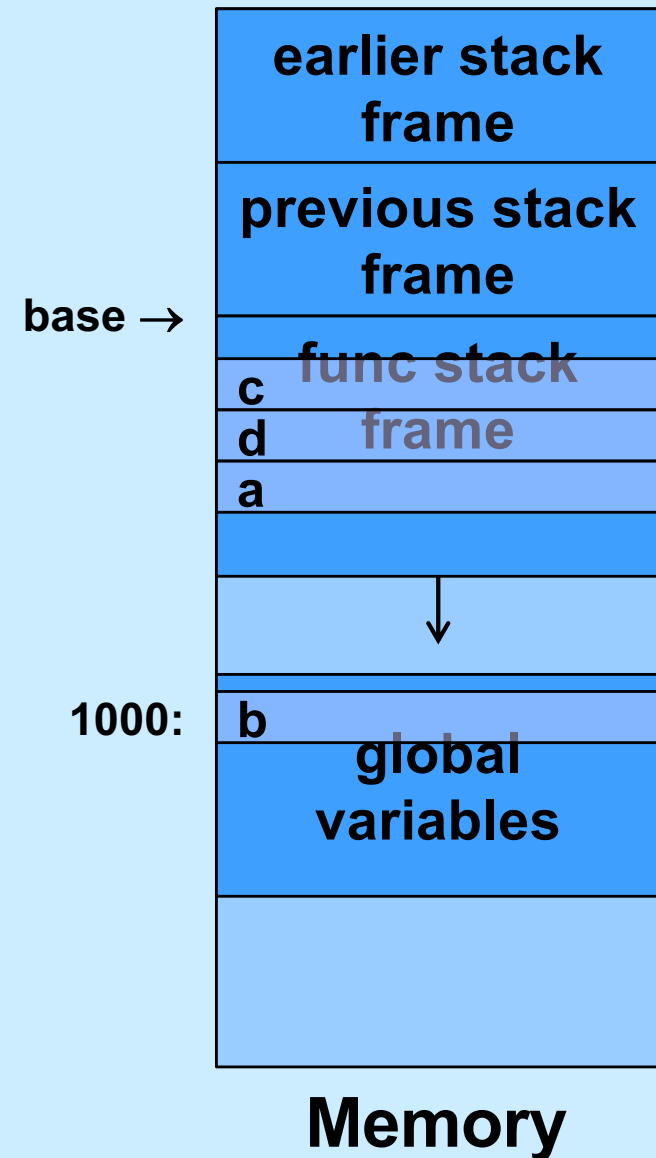


Quiz 4

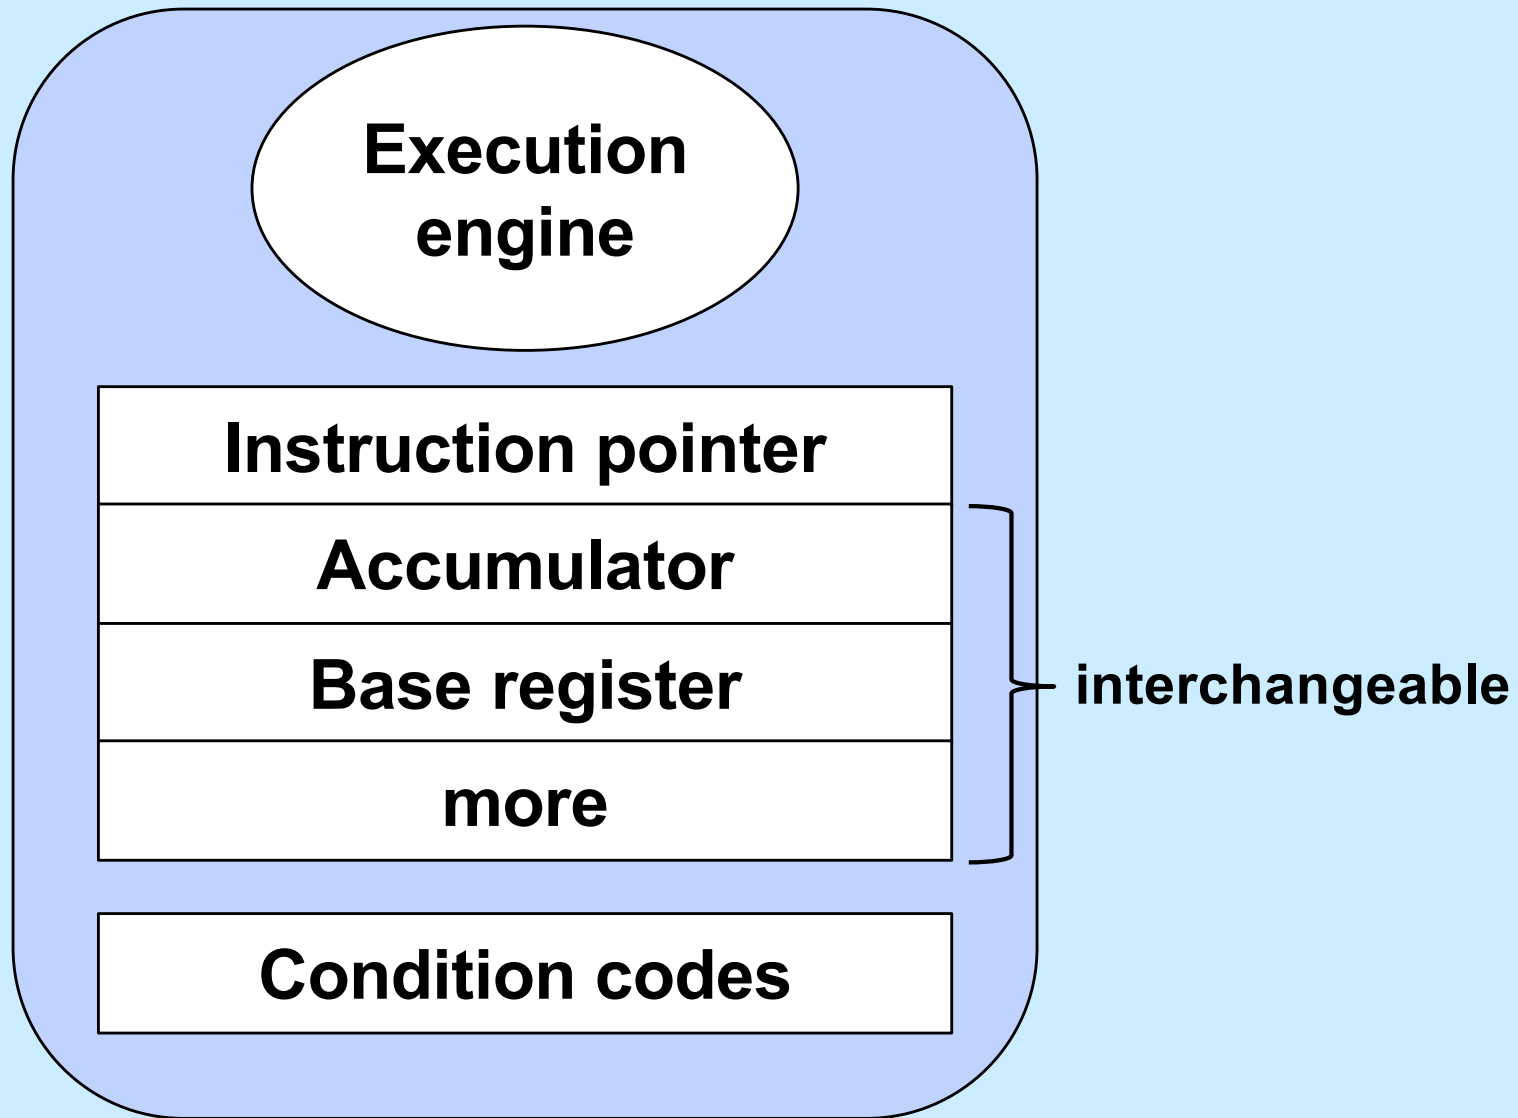
Suppose the value in *base* is 10,000. What is the address of *c*?

- a) 10,016
- b) 10,008
- c) 9992
- d) 9984

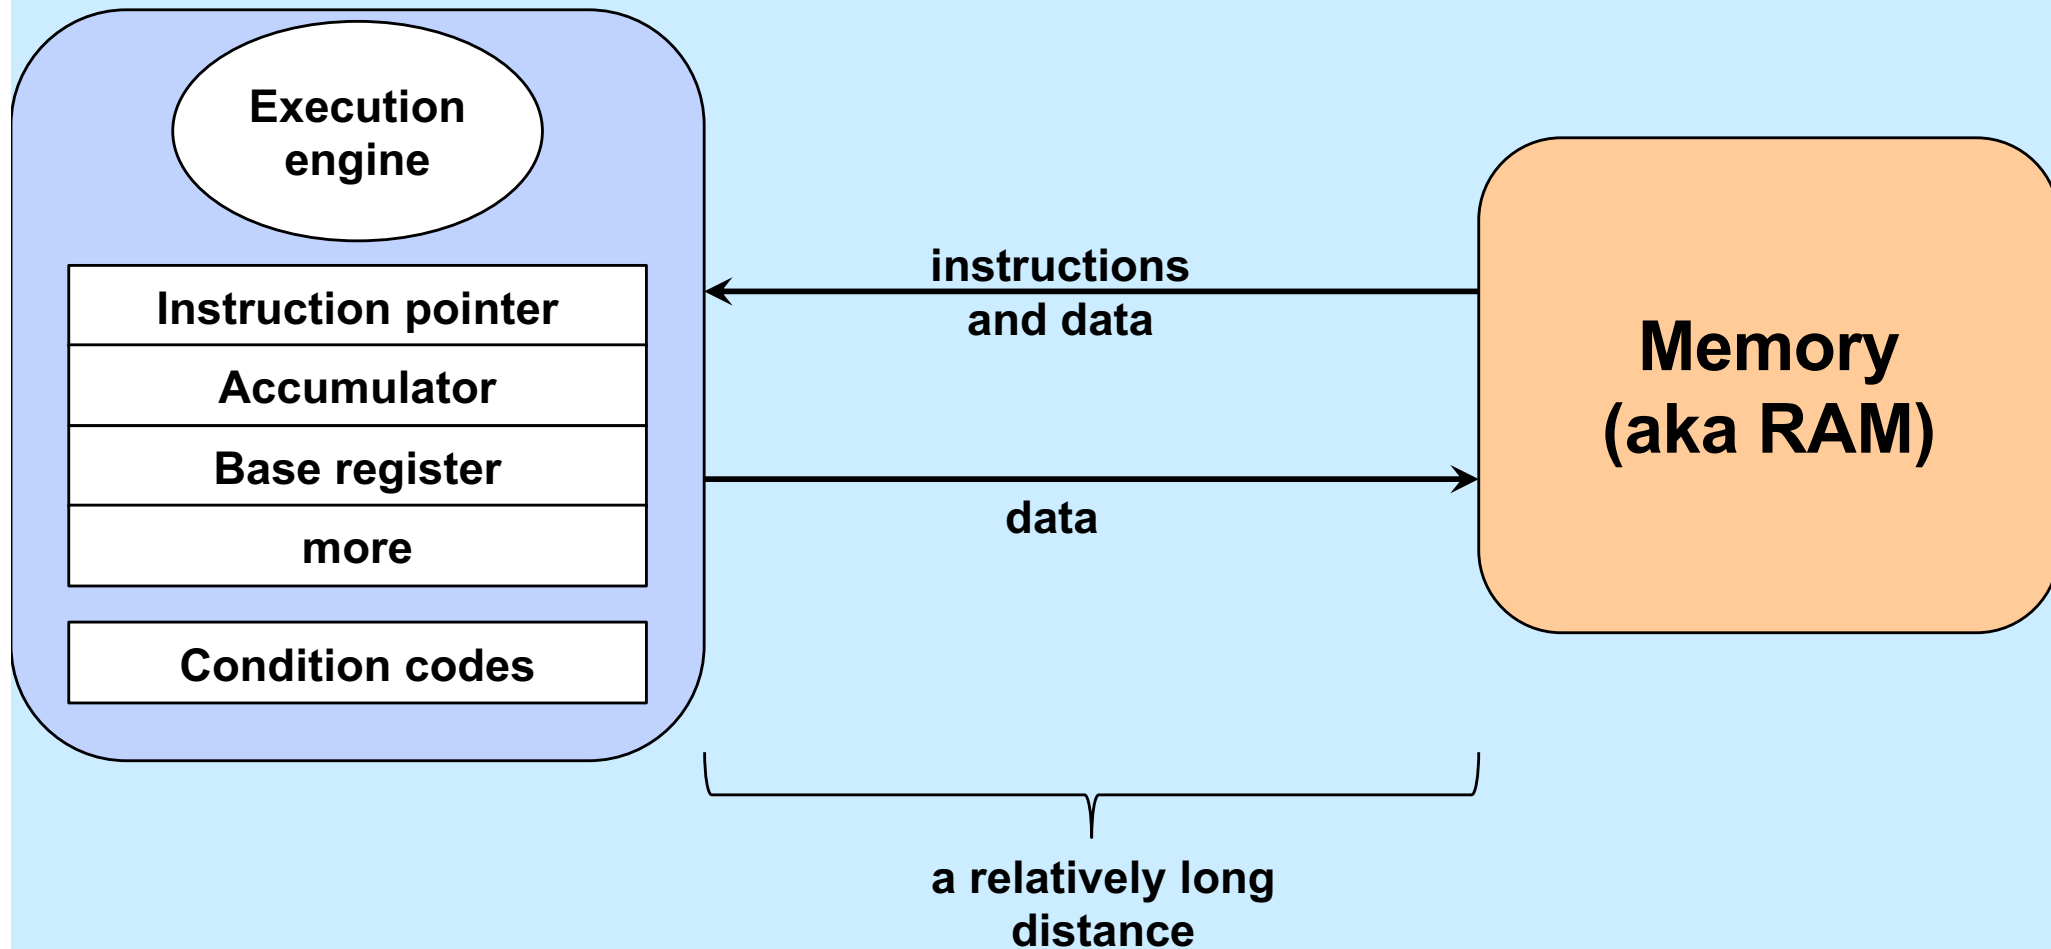
```
mov    1000, %acc
add    -8(%base), %acc
mul    -12(%base), %acc
mov    %acc, -16(%base)
```



Registers



Registers vs. Memory



Intel x86

- Intel created the 8008 (in 1972)
- 8008 begat 8080
- 8080 begat 8086
- 8086 begat 8088
- 8086 begat 286
- 286 begat 386
- 386 begat 486
- 486 begat Pentium
- Pentium begat Pentium Pro
- Pentium Pro begat Pentium II
- ad infinitum



IA32

2^{64}

- **2^{32} used to be considered a large number**
 - one couldn't afford 2^{32} bytes of memory, so no problem with that as an upper bound
- **Intel (and others) saw need for machines with 64-bit addresses**
 - devised IA64 architecture with HP
 - » became known as Itanium
 - » very different from x86
- **AMD also saw such a need**
 - developed 64-bit extension to x86, called x86-64
- **Itanium flopped**
- **x86-64 dominated**
- **Intel, reluctantly, adopted x86-64**

Why Intel?

- **Most CS Department machines are Intel**
- **An increasing number of personal machines are not**
 - **Apple has switched to ARM**
 - **packaged into their M1, M2, etc. chips**
 - » **“Apple Silicon”**
- **Intel x86-64 is very different from ARM64 — internally**
- **Programming concepts are similar**
- **We cover Intel; most of the concepts apply to ARM**