


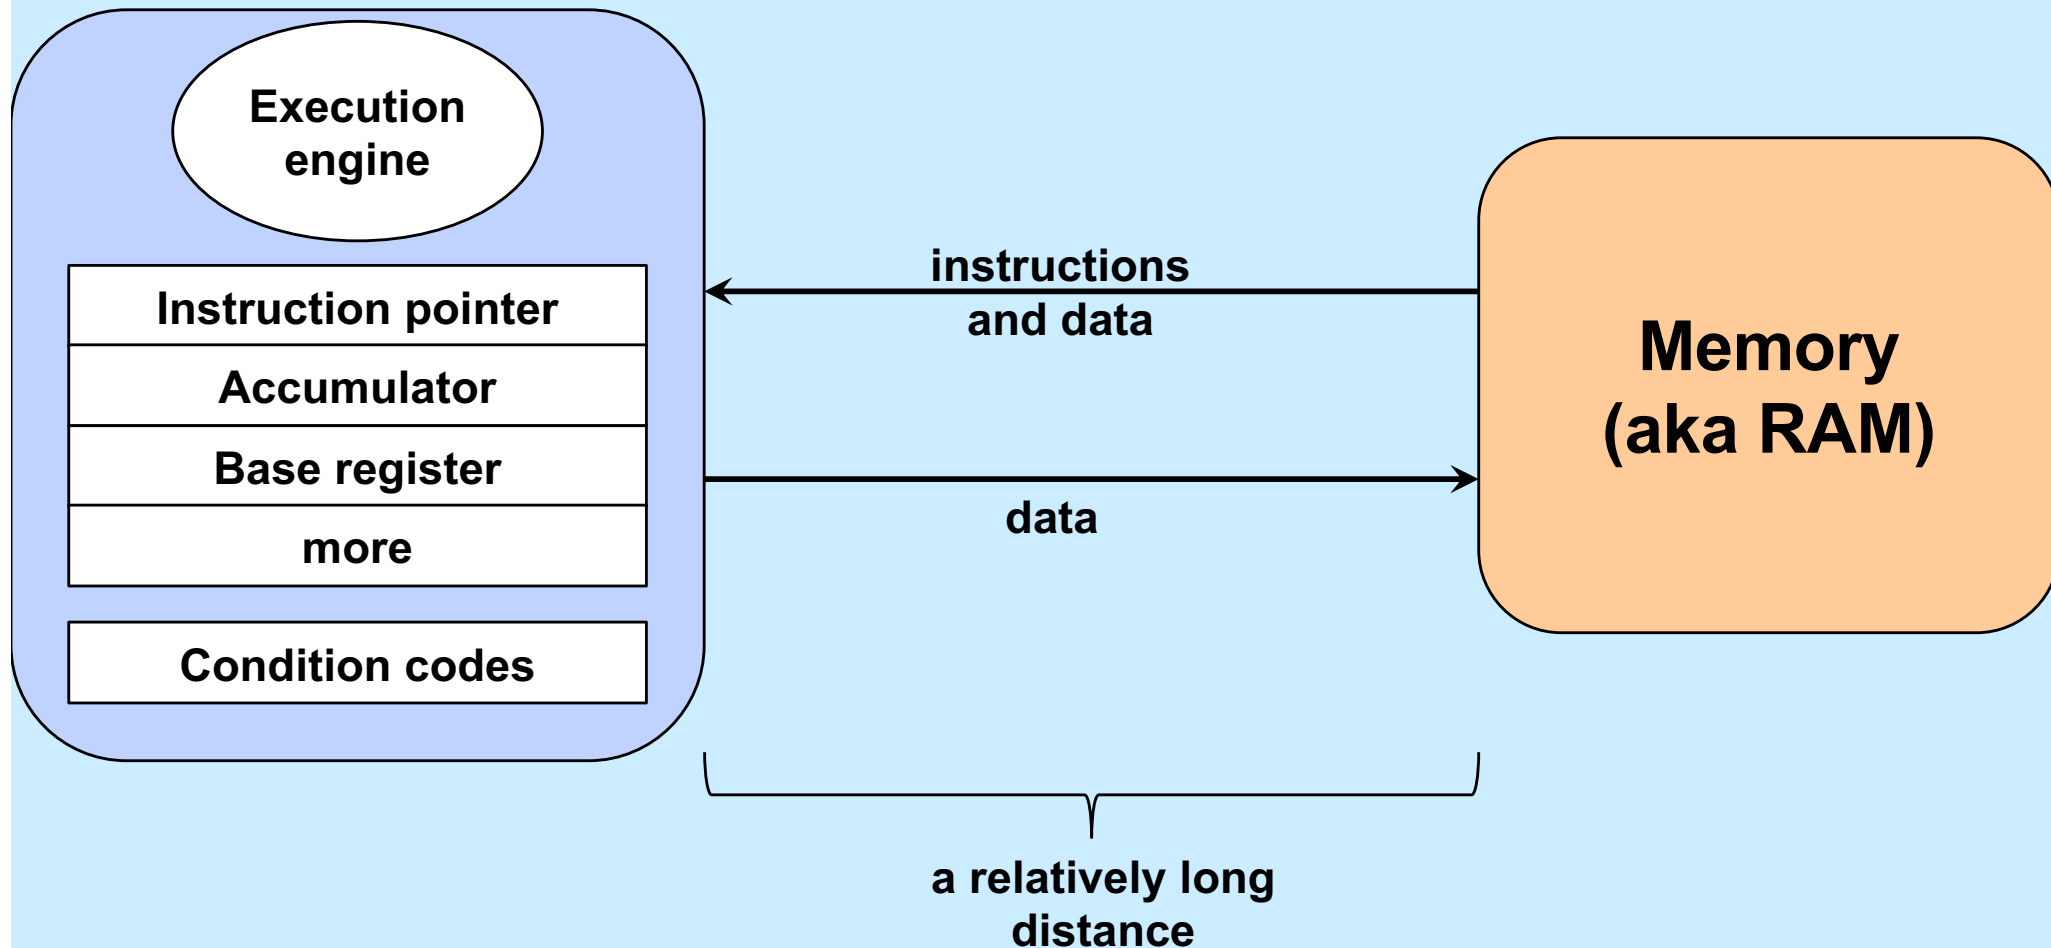
CS 33

Machine Programming (2)

Registers



Registers vs. Memory



Intel x86

- Intel created the 8008 (in 1972)
- 8008 begat 8080
- 8080 begat 8086
- 8086 begat 8088
- 8086 begat 286
- 286 begat 386
- 386 begat 486
- 486 begat Pentium
- Pentium begat Pentium Pro
- Pentium Pro begat Pentium II
- ad infinitum



IA32


2^{64}

- **2^{32} used to be considered a large number**
 - one couldn't afford 2^{32} bytes of memory, so no problem with that as an upper bound
- **Intel (and others) saw need for machines with 64-bit addresses**
 - devised IA64 architecture with HP
 - » became known as Itanium
 - » very different from x86
- **AMD also saw such a need**
 - developed 64-bit extension to x86, called x86-64
- **Itanium flopped**
- **x86-64 dominated**
- **Intel, reluctantly, adopted x86-64**

Data Types on IA32 and x86-64

- **“Integer” data of 1, 2, or 4 bytes (plus 8 bytes on x86-64)**
 - data values
 - » whether signed or unsigned depends on interpretation
 - addresses (untyped pointers)
- **Floating-point data of 4, 8, or 10 bytes**
- **No aggregate types such as arrays or structures**
 - just contiguously allocated bytes in memory

Operand Size



- Rather than `mov ...`
 - `movb`
 - `movs`
 - `movl`
 - `movq` (x86-64 only)

General-Purpose Registers (IA32)

Origin
(mostly obsolete)

general purpose	%eax	%ax	%ah	%al	accumulate
	%ecx	%cx	%ch	%cl	counter
	%edx	%dx	%dh	%dl	data
	%ebx	%bx	%bh	%bl	base
	%esi	%si			source index
	%edi	%di			destination index
	%esp	%sp			stack pointer
	%ebp	%bp			base pointer
16-bit virtual registers (backwards compatibility)					

x86-64 General-Purpose Registers

	%rax	%eax	%r8	%r8d	a5
	%rbx	%ebx	%r9	%r9d	a6
a4	%rcx	%ecx	%r10	%r10d	
a3	%rdx	%edx	%r11	%r11d	
a2	%rsi	%esi	%r12	%r12d	
a1	%rdi	%edi	%r13	%r13d	
	%rsp	%esp	%r14	%r14d	
	%rbp	%ebp	%r15	%r15d	

- Extend existing registers to 64 bits. Add 8 new ones.

Moving Data

- Moving data

`movq source, dest`

- Operand types

- **Immediate:** constant integer data

- » example: `$0x400`, `$-533`

- » like C constant, but prefixed with ``$'`

- » encoded with 1, 2, 4, or 8 bytes

- **Register:** one of 16 64-bit registers

- » example: `%rax`, `%rdx`

- » `%rsp` and `%rbp` have some special uses

- » others have special uses for particular instructions

- **Memory:** 8 consecutive bytes of memory at address given by register(s)

- » simplest example: `(%rax)`

- » various other “address modes”

`%rax`

`%rcx`

`%rdx`

`%rbx`

`%rsi`

`%rdi`

`%rsp`

`%rbp`

`%r8`

`%r9`

`%r10`

`%r11`

`%r12`

`%r13`

`%r14`

`%r15`

movq Operand Combinations

	Source	Dest	Src, Dest	C Analog
movq	Imm	Reg	movq \$0x4,%rax	temp = 0x4;
		Mem	movq \$-147, (%rax)	*p = -147;
	Reg	Reg	movq %rax,%rdx	temp2 = temp1;
		Mem	movq %rax, (%rdx)	*p = temp;
	Mem	Reg	movq (%rax), %rdx	temp = *p;

Cannot (normally) do memory-memory transfer with a single instruction

Simple Memory Addressing Modes

- Normal (R) Mem[Reg[R]]
 - register R specifies memory address

```
movq (%rcx), %rax
```

- Displacement D(R) Mem[Reg[R]+D]
 - register R specifies start of memory region
 - constant displacement D specifies offset

```
movq 8(%rbp), %rdx
```

Using Simple Addressing Modes


```
struct xy {  
    long x;  
    long y;  
}  
  
void swapxy(struct xy *p) {  
    long temp = p->x;  
    p->x = p->y;  
    p->y = temp;  
}
```

swap:

```
movq (%rdi), %rax  
movq 8(%rdi), %rdx  
movq %rdx, (%rdi)  
movq %rax, 8(%rdi)  
ret
```

Understanding Swapxy

```
struct xy {  
    long x;  
    long y;  
}  
  
void swapxy(struct xy *p) {  
    long temp = p->x;  
    p->x = p->y;  
    p->y = temp;  
}
```

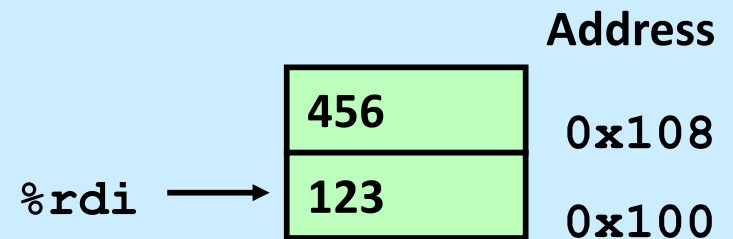


Register	Value
%rdi	p
%rax	temp
%rdx	p->y

```
movq (%rdi), %rax    # temp = p->x  
movq 8(%rdi), %rdx   # %rdx = p->y  
movq %rdx, (%rdi)    # p->x = %rdx  
movq %rax, 8(%rdi)   # p->y = temp  
ret
```

Understanding Swapxy


<code>%rdi</code>	<code>0x100</code>
<code>%rax</code>	
<code>%rdx</code>	



```
movq (%rdi), %rax      # temp = p->x
movq 8(%rdi), %rdx     # %rdx = p->y
movq %rdx, (%rdi)      # p->x = %rdx
movq %rax, 8(%rdi)     # p->y = temp
ret
```

Understanding Swapxy


%rdi	0x100
%rax	123
%rdx	



```
movq (%rdi), %rax      # temp = p->x
movq 8(%rdi), %rdx     # %rdx = p->y
movq %rdx, (%rdi)      # p->x = %rdx
movq %rax, 8(%rdi)     # p->y = temp
ret
```


Understanding Swapxy


%rdi	0x100
%rax	123
%rdx	456



```
movq (%rdi), %rax      # temp = p->x
movq 8(%rdi), %rdx     # %rdx = p->y
movq %rdx, (%rdi)      # p->x = %rdx
movq %rax, 8(%rdi)     # p->y = temp
ret
```

Understanding Swapxy


%rdi	0x100
%rax	123
%rdx	456



```
movq (%rdi), %rax      # temp = p->x
movq 8(%rdi), %rdx     # %rdx = p->y
movq %rdx, (%rdi)      # p->x = %rdx
movq %rax, 8(%rdi)     # p->y = temp
ret
```

Understanding Swapxy


%rdi	0x100
%rax	123
%rdx	456



```
movq (%rdi), %rax      # temp = p->x
movq 8(%rdi), %rdx     # %rdx = p->y
movq %rdx, (%rdi)      # p->x = %rdx
movq %rax, 8(%rdi)     # p->y = temp
ret
```

Quiz 1

```
movq -8(%rbp), %rax
movq (%rax), %rax
movq (%rax), %rax
movq %rax, -16(%rbp)
```



Which C statements best describe the assembler code?

```
// a
long x;
long y;
y = x;
```

```
// b
long *x;
long y;
y = *x;
```

```
// c
long **x;
long y;
y = **x;
```

```
// d
long ***x;
long y;
y = ***x;
```

Complete Memory-Addressing Modes

- Most general form

$D(Rb, Ri, S)$ $Mem[Reg[Rb] + S * Reg[Ri] + D]$

- D: constant “displacement”
- Rb: base register: any of 16[†] registers
- Ri: index register: any, except for `%rsp`
- S: scale: 1, 2, 4, or 8

- Special cases

(Rb, Ri)	$Mem[Reg[Rb] + Reg[Ri]]$
$D(Rb, Ri)$	$Mem[Reg[Rb] + Reg[Ri] + D]$
(Rb, Ri, S)	$Mem[Reg[Rb] + S * Reg[Ri]]$
D	$Mem[D]$

[†]The instruction pointer may also be used (for a total of 17 registers)

Address-Computation Examples

%rdx	0xf000
%rcx	0x0100

Expression	Address Computation	Address
0x8(%rdx)	0xf000 + 0x8	0xf008
(%rdx, %rcx)	0xf000 + 0x100	0xf100
(%rdx, %rcx, 4)	0xf000 + 4*0x0100	0xf400
0x80(,%rdx, 2)	2*0xf000 + 0x80	0x1e080

Address-Computation Instruction

- **leaq src, dest**
 - src is address mode expression
 - set *dest* to address denoted by expression
- **Uses**
 - computing addresses without a memory reference
 - » e.g., translation of `p = &x[i];`
 - computing arithmetic expressions of the form $x + k*y$
 - » $k = 1, 2, 4, \text{ or } 8$
- **Example**

```
long mul12(long x)
{
    return x*12;
}
```

Converted to ASM by compiler:

```
                                # x is in %rdi
leaq (%rdi,%rdi,2), %rax        # t <- x+x*2
shlq $2, %rax                  # return t<<2
```

32-bit Operands on x86-64

- **addl 4(%rdx), %eax**
 - memory address must be 64 bits
 - operands (in this case) are 32-bit
 - » result goes into %eax
 - lower half of %rax
 - upper half is filled with zeroes

Quiz 2


What value ends up in %ecx?

```
movq $1000, %rax
movq $1, %rbx
movl 2(%rax, %rbx, 2), %ecx
```

- a) 0x04050607
- b) 0x07060504
- c) 0x06070809
- d) 0x09080706

1009:	0x09
1008:	0x08
1007:	0x07
1006:	0x06
1005:	0x05
1004:	0x04
1003:	0x03
1002:	0x02
1001:	0x01
%rax → 1000:	0x00

Hint:



Swapxy for Ints

```
struct xy {  
    int x;  
    int y;  
}  
  
void swapxy(struct xy *p) {  
    int temp = p->x;  
    p->x = p->y;  
    p->y = temp;  
}
```

swap:

```
movl (%rdi), %eax  
movl 4(%rdi), %edx  
movl %edx, (%rdi)  
movl %eax, 4(%rdi)  
ret
```

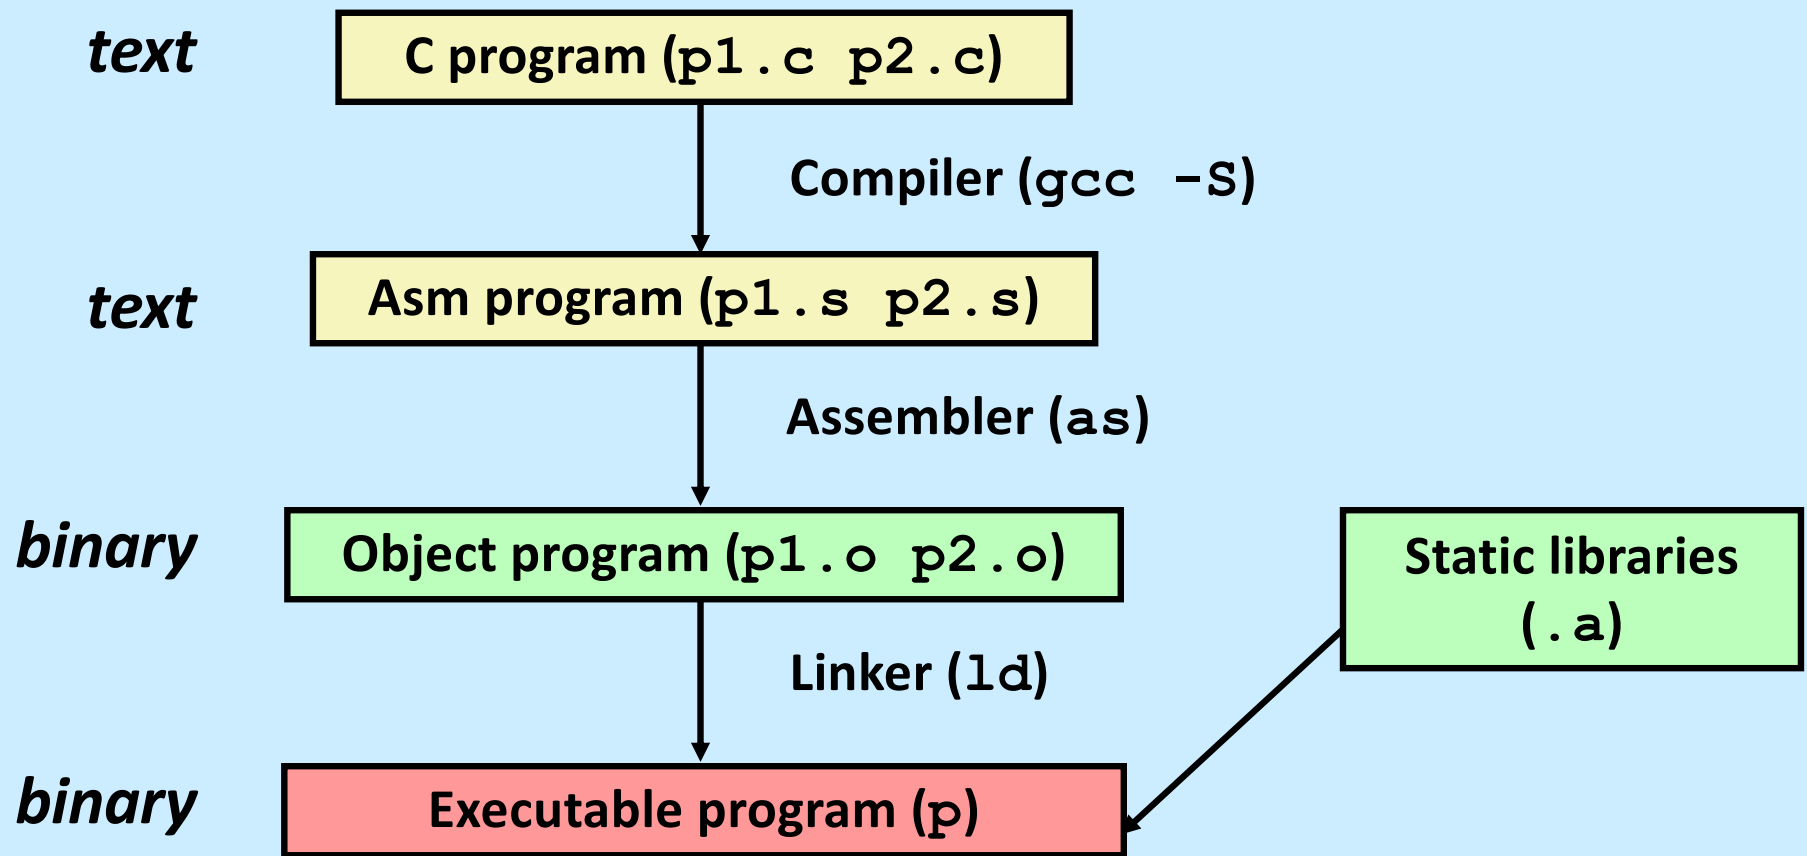
- **Pointers are 64 bits**
- **What they point to are 32 bits**

Bytes

- **Each register has a byte version**
 - e.g., %r10: %r10b; see earlier slide for x86 registers
- **Needed for byte instructions**
 - `movb (%rax, %rsi), %r10b`
 - sets *only* the low byte in %r10
 - » other seven bytes are unchanged
- **Alternatives**
 - `movzbq (%rax, %rsi), %r10`
 - » copies byte to low byte of %r10
 - » zeroes go to higher bytes
 - `movsbq (%rax, %rsi), %r10`
 - » copies byte to low byte of %r10
 - » sign is extended to all higher bits

Turning C into Object Code

- Code in files `p1.c` `p2.c`
- Compile with command: `gcc -O1 p1.c p2.c -o p`
 - » use basic optimizations (`-O1`)
 - » put resulting binary in file `p`



Example

```
long ASum(long *a, unsigned long size) {  
    long i, sum = 0;  
    for (i=0; i<size; i++)  
        sum += a[i];  
    return sum;  
}
```

Object Code

Code for ASum

0x112b <ASum>:

0x48

0x85

0xf6

0x74

0x19

0x48

0x89

0xfa

0x48

0x8d

0x0c

0xf7

.

.

.

- Total of 35 bytes

- Each instruction:
1, 2, or 3 bytes

- Starts at address
0x112b

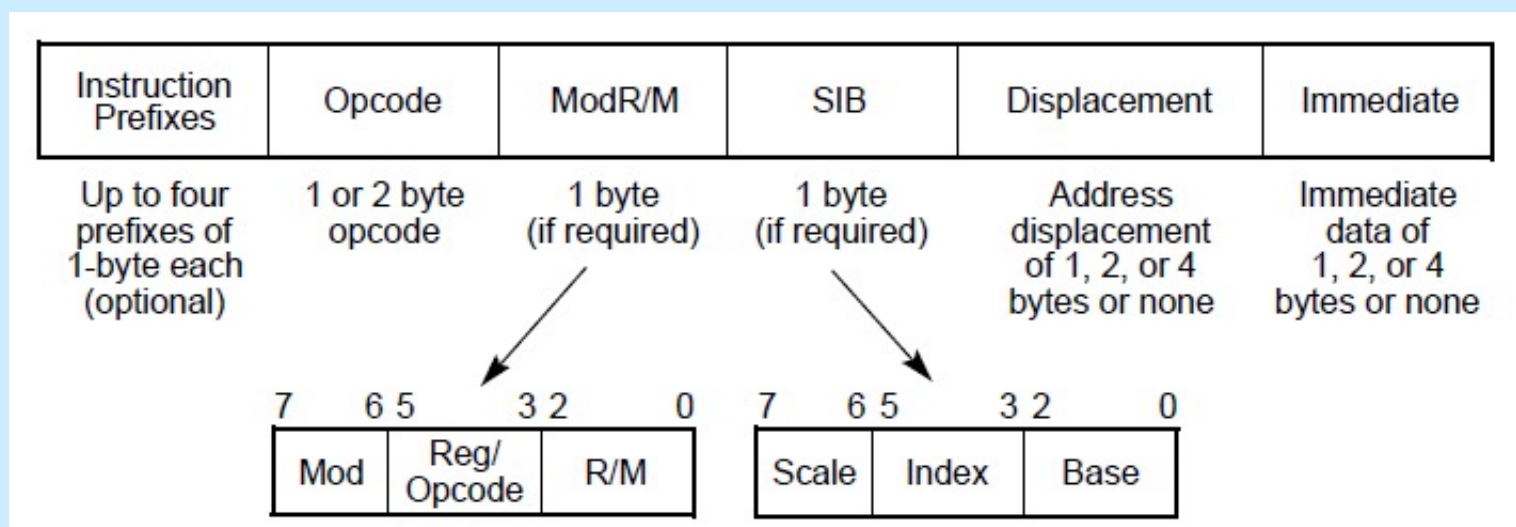
- **Assembler**

- translates .s into .o
- binary encoding of each instruction
- nearly-complete image of executable code
- missing linkages between code in different files

- **Linker**

- resolves references between files
- combines with static run-time libraries
 - » e.g., code for printf
- some libraries are *dynamically linked*
 - » linking occurs when program begins execution

Instruction Format



Disassembling Object Code

Disassembled

```
0000000000000112b <ASum>:
   112b:  48 85 f6          test    %rsi,%rsi
   112e:  74 19             je      1149 <ASum+0x1e>
   1130:  48 89 fa          mov     %rdi,%rdx
   1133:  48 8d 0c f7        lea     (%rdi,%rsi,8),%rcx
   1137:  b8 00 00 00 00    mov     $0x0,%eax
   113c:  48 03 02          add     (%rdx),%rax
   113f:  48 83 c2 08        add     $0x8,%rdx
   1143:  48 39 ca          cmp     %rcx,%rdx
   1146:  75 f4             jne     113c <ASum+0x11>
   1148:  c3               retq
   1149:  b8 00 00 00 00    mov     $0x0,%eax
   114e:  c3               retq
```

- **Disassembler**

`objdump -d <file>`

- useful tool for examining object code
- produces approximate rendition of assembly code

Alternate Disassembly

Object

0x112b:

0x48

0x85

0xf6

0x74

0x19

0x48

0x89

0xfa

0x48

0x8d

0x0c

0xf7

.

.

.

Disassembled

Dump of assembler code for function ASum:

0x112b <+0>: test %rsi,%rsi

0x112e <+3>: je 0x1149 <ASum+30>

0x1130 <+5>: mov %rdi,%rdx

0x1133 <+8>: lea (%rdi,%rsi,8),%rcx

0x1137 <+12>: mov \$0x0,%eax

...

- **Within gdb debugger**

`gdb <file>`

`disassemble ASum`

– disassemble the ASum object code

`x/35xb ASum`

– examine the 35 bytes starting at ASum

How Many Instructions are There?

- We cover ~30
- Implemented by Intel:
 - 80 in original 8086 architecture
 - 7 added with 80186
 - 17 added with 80286
 - 33 added with 386
 - 6 added with 486
 - 6 added with Pentium
 - 1 added with Pentium MMX
 - 4 added with Pentium Pro
 - 8 added with SSE
 - 8 added with SSE2
 - 2 added with SSE3
 - 14 added with x86-64
 - 10 added with VT-x
 - 2 added with SSE4a
- Total: 198
- Doesn't count:
 - floating-point instructions
 - » ~100
 - SIMD instructions
 - » lots
 - AMD-added instructions
 - undocumented instructions

Some Arithmetic Operations

- Two-operand instructions:

Format	Computation	
<code>addl</code>	<code>Src, Dest</code>	<code>Dest = Dest + Src</code>
<code>subl</code>	<code>Src, Dest</code>	<code>Dest = Dest - Src</code>
<code>imull</code>	<code>Src, Dest</code>	<code>Dest = Dest * Src</code>
<code>shll</code>	<code>Src, Dest</code>	<code>Dest = Dest << Src</code>
<code>sarl</code>	<code>Src, Dest</code>	<code>Dest = Dest >> Src</code>
<code>shrl</code>	<code>Src, Dest</code>	<code>Dest = Dest >> Src</code>
<code>xorl</code>	<code>Src, Dest</code>	<code>Dest = Dest ^ Src</code>
<code>andl</code>	<code>Src, Dest</code>	<code>Dest = Dest & Src</code>
<code>orl</code>	<code>Src, Dest</code>	<code>Dest = Dest Src</code>

Also called `sall`
Arithmetic
Logical

– watch out for argument order!

Some Arithmetic Operations

- **One-operand Instructions**

`incl` `Dest` $= \text{Dest} + 1$

`decl` `Dest` $= \text{Dest} - 1$

`negl` `Dest` $= -\text{Dest}$

`notl` `Dest` $= \sim\text{Dest}$

- **See textbook for more instructions**
- **See Intel documentation for even more**

Arithmetic Expression Example

```
int arith(int x, int y, int z)
{
    int t1 = x+y;
    int t2 = z+t1;
    int t3 = x+4;
    int t4 = y * 48;
    int t5 = t3 + t4;
    int rval = t2 * t5;
    return rval;
}
```

```
arith:
    leal    (%rdi,%rsi), %eax
    addl    %edx, %eax
    leal    (%rsi,%rsi,2), %edx
    shll    $4, %edx
    leal    4(%rdi,%rdx), %ecx
    imull    %ecx, %eax
    ret
```

Understanding arith

```
int arith(int x, int y, int z)
{
    int t1 = x+y;
    int t2 = z+t1;
    int t3 = x+4;
    int t4 = y * 48;
    int t5 = t3 + t4;
    int rval = t2 * t5;
    return rval;
}
```

```
leal    (%rdi,%rsi), %eax
addl    %edx, %eax
leal    (%rsi,%rsi,2), %edx
shll    $4, %edx
leal    4(%rdi,%rdx), %ecx
imull    %ecx, %eax
ret
```

%rdx	z
%rsi	y
%rdi	x

Understanding arith

```
int arith(int x, int y, int z)
{
    int t1 = x+y;
    int t2 = z+t1;
    int t3 = x+4;
    int t4 = y * 48;
    int t5 = t3 + t4;
    int rval = t2 * t5;
    return rval;
}
```

%rdx	z
%rsi	y
%rdi	x

```
leal    (%rdi,%rsi), %eax    # eax = x+y      (t1)
addl    %edx, %eax          # eax = t1+z      (t2)
leal    (%rsi,%rsi,2), %edx  # edx = 3*y      (t4)
shll    $4, %edx            # edx = t4*16     (t4)
leal    4(%rdi,%rdx), %ecx   # ecx = x+4+t4   (t5)
imull   %ecx, %eax          # eax *= t5        (rval)
ret
```

Observations about `arith`

```
int arith(int x, int y, int z)
{
    int t1 = x+y;
    int t2 = z+t1;
    int t3 = x+4;
    int t4 = y * 48;
    int t5 = t3 + t4;
    int rval = t2 * t5;
    return rval;
}
```

- Instructions in different order from C code
- Some expressions might require multiple instructions
- Some instructions might cover multiple expressions

<code>leal</code>	<code>(%rdi,%rsi), %eax</code>	<code>#</code>	<code>eax = x+y</code>	<code>(t1)</code>
<code>addl</code>	<code>%edx, %eax</code>	<code>#</code>	<code>eax = t1+z</code>	<code>(t2)</code>
<code>leal</code>	<code>(%rsi,%rsi,2), %edx</code>	<code>#</code>	<code>edx = 3*y</code>	<code>(t4)</code>
<code>shll</code>	<code>\$4, %edx</code>	<code>#</code>	<code>edx = t4*16</code>	<code>(t4)</code>
<code>leal</code>	<code>4(%rdi,%rdx), %ecx</code>	<code>#</code>	<code>ecx = x+4+t4</code>	<code>(t5)</code>
<code>imull</code>	<code>%ecx, %eax</code>	<code>#</code>	<code>eax *= t5</code>	<code>(rval)</code>
<code>ret</code>				

Another Example

```
int logical(int x, int y)
{
    int t1 = x^y;
    int t2 = t1 >> 17;
    int mask = (1<<13) - 7;
    int rval = t2 & mask;
    return rval;
}
```

$2^{13} = 8192, 2^{13} - 7 = 8185$

<code>xorl %esi, %edi</code>	<code># edi = x^y</code>	<code>(t1)</code>
<code>sarl \$17, %edi</code>	<code># edi = t1>>17</code>	<code>(t2)</code>
<code>movl %edi, %eax</code>	<code># eax = edi</code>	
<code>andl \$8185, %eax</code>	<code># eax = t2 & mask</code>	<code>(rval)</code>

Quiz 3

- What is the final value in %ecx?

```
xorl %ecx, %ecx
```

```
incl %ecx
```

```
shll %cl, %ecx # %cl is the low byte of %ecx
```

```
addl %ecx, %ecx
```

a) 0

b) 2

c) 4

d) 8