

# CS 33

## Files Part 3

# Setting File Permissions

```
#include <sys/types.h>
#include <sys/stat.h>
int chmod(const char *path, mode_t mode)
```

- sets the file permissions of the given file to those specified in *mode*
- only the owner of a file and the superuser may change its permissions
- nine combinable possibilities for *mode* (*read/write/execute* for *user*, *group*, and *others*)
  - » S\_IRUSR (0400), S\_IWUSR (0200), S\_IXUSR (0100)
  - » S\_IRGRP (040), S\_IWGRP (020), S\_IXGRP (010)
  - » S\_IROTH (04), S\_IWOTH (02), S\_IXOTH (01)

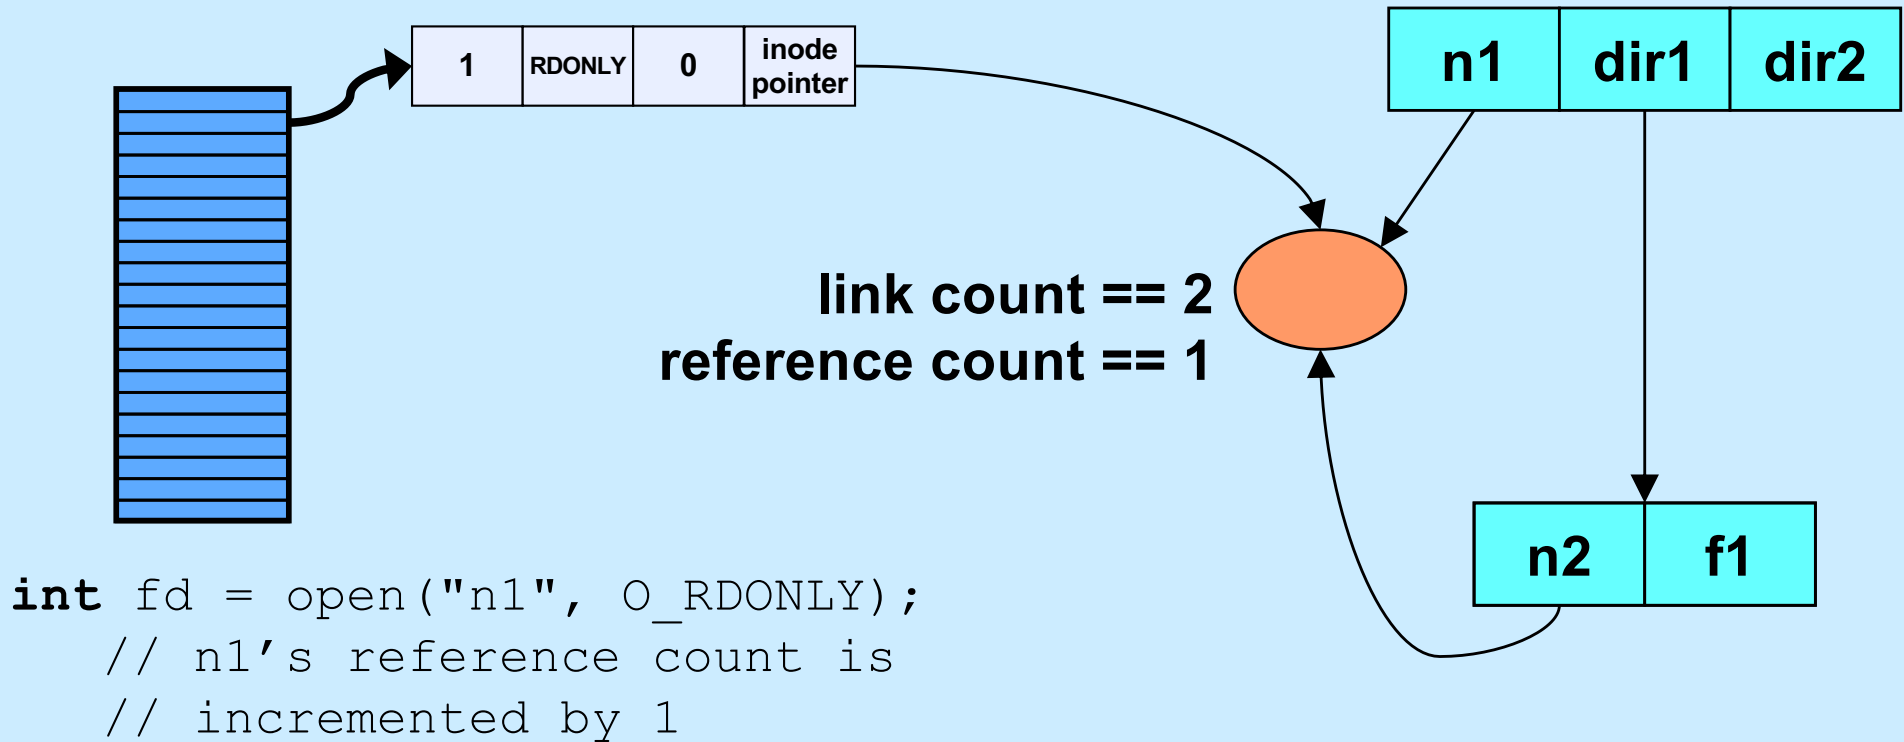
# Umask

- **Standard programs create files with “maximum needed permissions” as mode**
  - compilers: 0777
  - editors: 0666
- **Per-process parameter, *umask*, used to turn off undesired permission bits**
  - e.g., turn off all permissions for others, write permission for group: set umask to 027
    - » compilers: permissions =  $0777 \& \sim(027) = 0750$
    - » editors: permissions =  $0666 \& \sim(027) = 0640$
  - set with *umask* system call or (usually) shell command

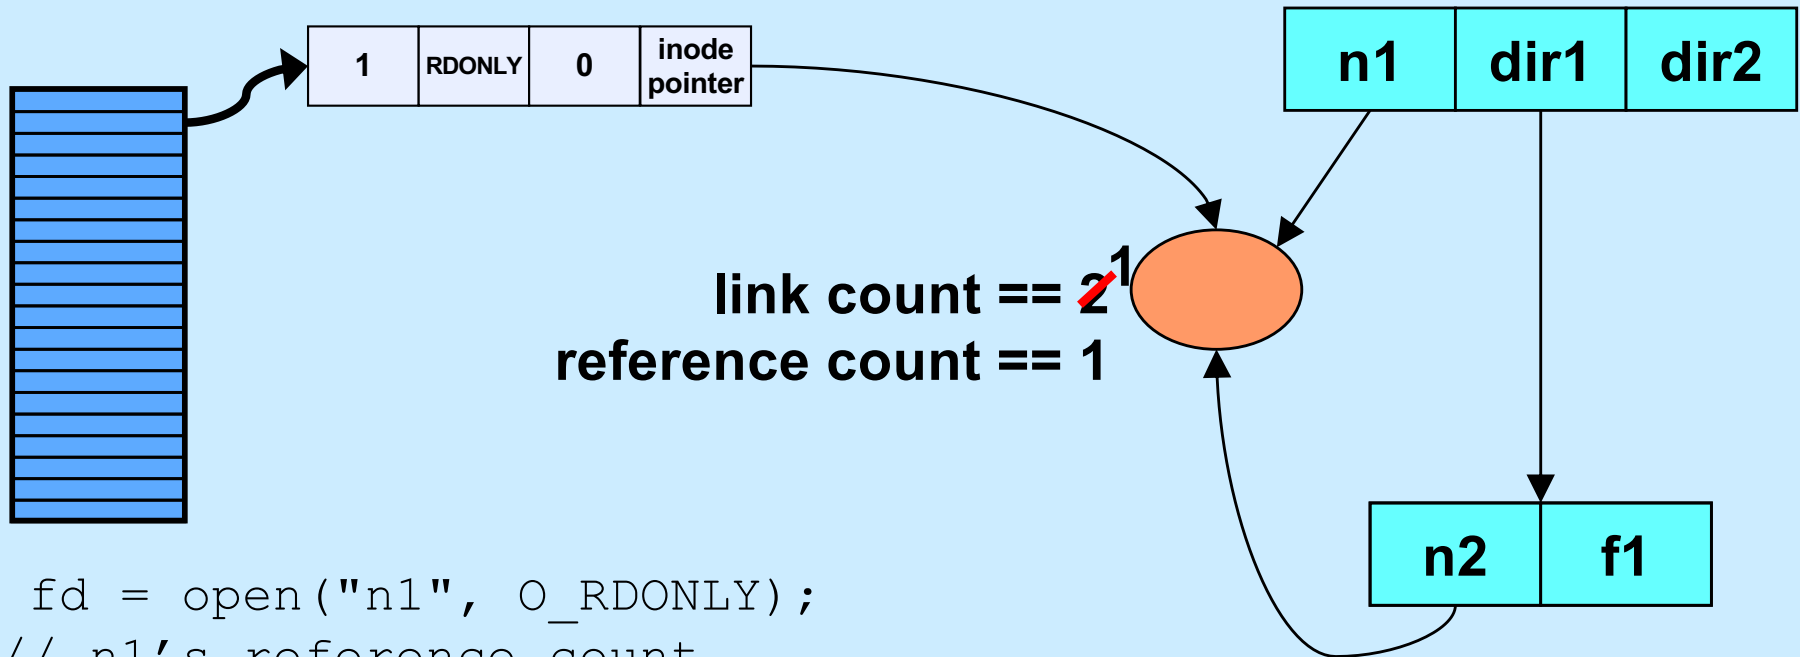
# Creating a File

- Use either *open* or *creat*
  - `open(const char *pathname, int flags, mode_t mode)`
    - » flags must include `O_CREAT`
  - `creat(const char *pathname, mode_t mode)`
    - » `open` is preferred
- The *mode* parameter helps specify the permissions of the newly created file
  - `permissions = mode & ~umask`

# Link and Reference Counts



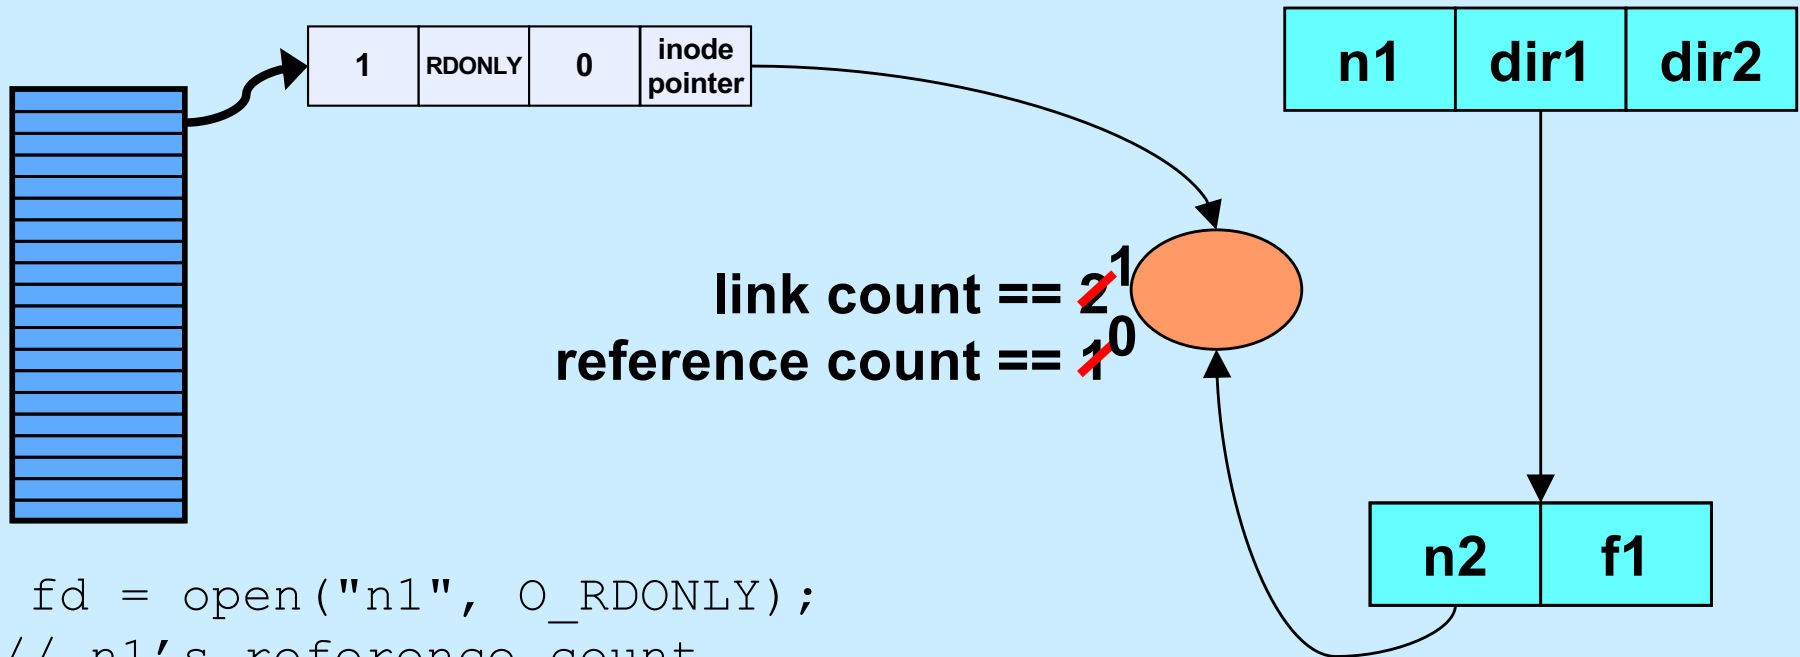
# Link and Reference Counts



```
int fd = open("n1", O_RDONLY);  
// n1's reference count  
// incremented by 1
```


```
unlink("n1");  
// link count decremented by 1  
// same effect in shell via "rm n1"
```

# Link and Reference Counts



```
int fd = open("n1", O_RDONLY);  
    // n1's reference count  
    // incremented by 1  
  
unlink("n1");  
    // link count decremented by 1  
  
close(fd);  
    // reference count decremented by 1
```

# Link and Reference Counts

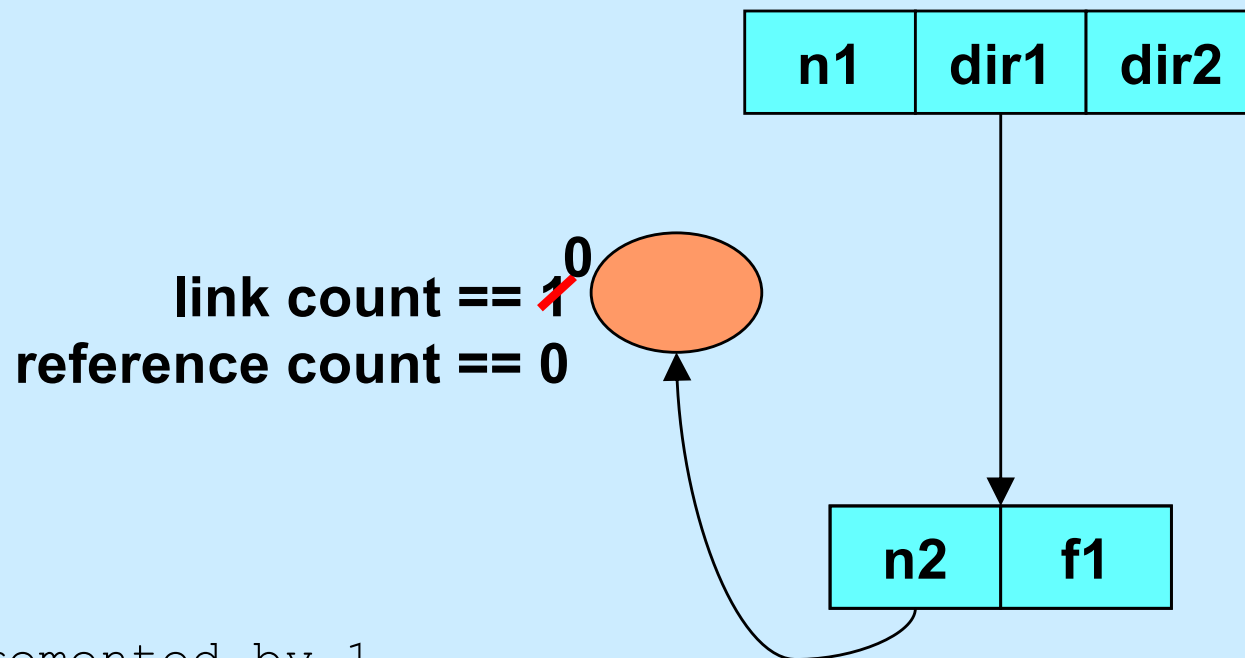


```
int fd = open("n1", O_RDONLY);  
    // n1's reference count  
    // incremented by 1  
  
unlink("n1");  
    // link count decremented by 1  
  
close(fd);  
    // reference count decremented by 1
```



# Link and Reference Counts

```
unlink("dir1/n2");  
// link count decremented by 1
```




# Quiz 1

```
int main() {  
    int fd = open("file", O_RDWR|O_CREAT, 0666);  
    unlink("file");  
    PutStuffInFile(fd);  
    GetStuffFromFile(fd);  
    return 0;  
}
```


**Assume that *PutStuffInFile* writes to the given file, and *GetStuffFromFile* reads from the file.**

- a) This program is doomed to failure, since the file is deleted before it's used**
- b) Because the file is used after the unlink call, it won't be deleted**
- c) The file will be deleted when the program terminates**


# Interprocess Communication (IPC): Pipes




# Interprocess Communication: Same Machine I



# Interprocess Communication: Same Machine II




# Interprocess Communication: Different Machines



# Pipes

```
$cs1ab2e who | wc -l
```



# Using Pipes in C

```
$cs1ab2e who | wc -l
```

```
int fd[2];
pipe(fd);
if (fork() == 0) {
    close(fd[0]);
    close(1);
    dup(fd[1]); close(fd[1]);
    execl("/usr/bin/who", "who", 0); // who sends output to pipe
}
if (fork() == 0) {
    close(fd[1]);
    close(0);
    dup(fd[0]); close(fd[0]);
    execl("/usr/bin/wc", "wc", "-l", 0); // wc's input is from pipe
}
close(fd[1]); close(fd[0]);
// ...
```






# Sharing Files

- **You're doing a project with a partner**
- **You code it as one 15,000-line file**
  - the first 7,500 lines are yours
  - the second 7,500 lines are your partner's
- **You edit the file, changing 6,000 lines**
  - it's now 5am
- **Your partner completes her changes at 5:01am**
- **At 5:02am you look at the file**
  - your partner's changes are there
  - yours are not

# Lessons

- **Never work with a partner**
- **Use more than one file**
- **Read up on git**
- **Use an editor and file system that support file locking**


# What We Want ...



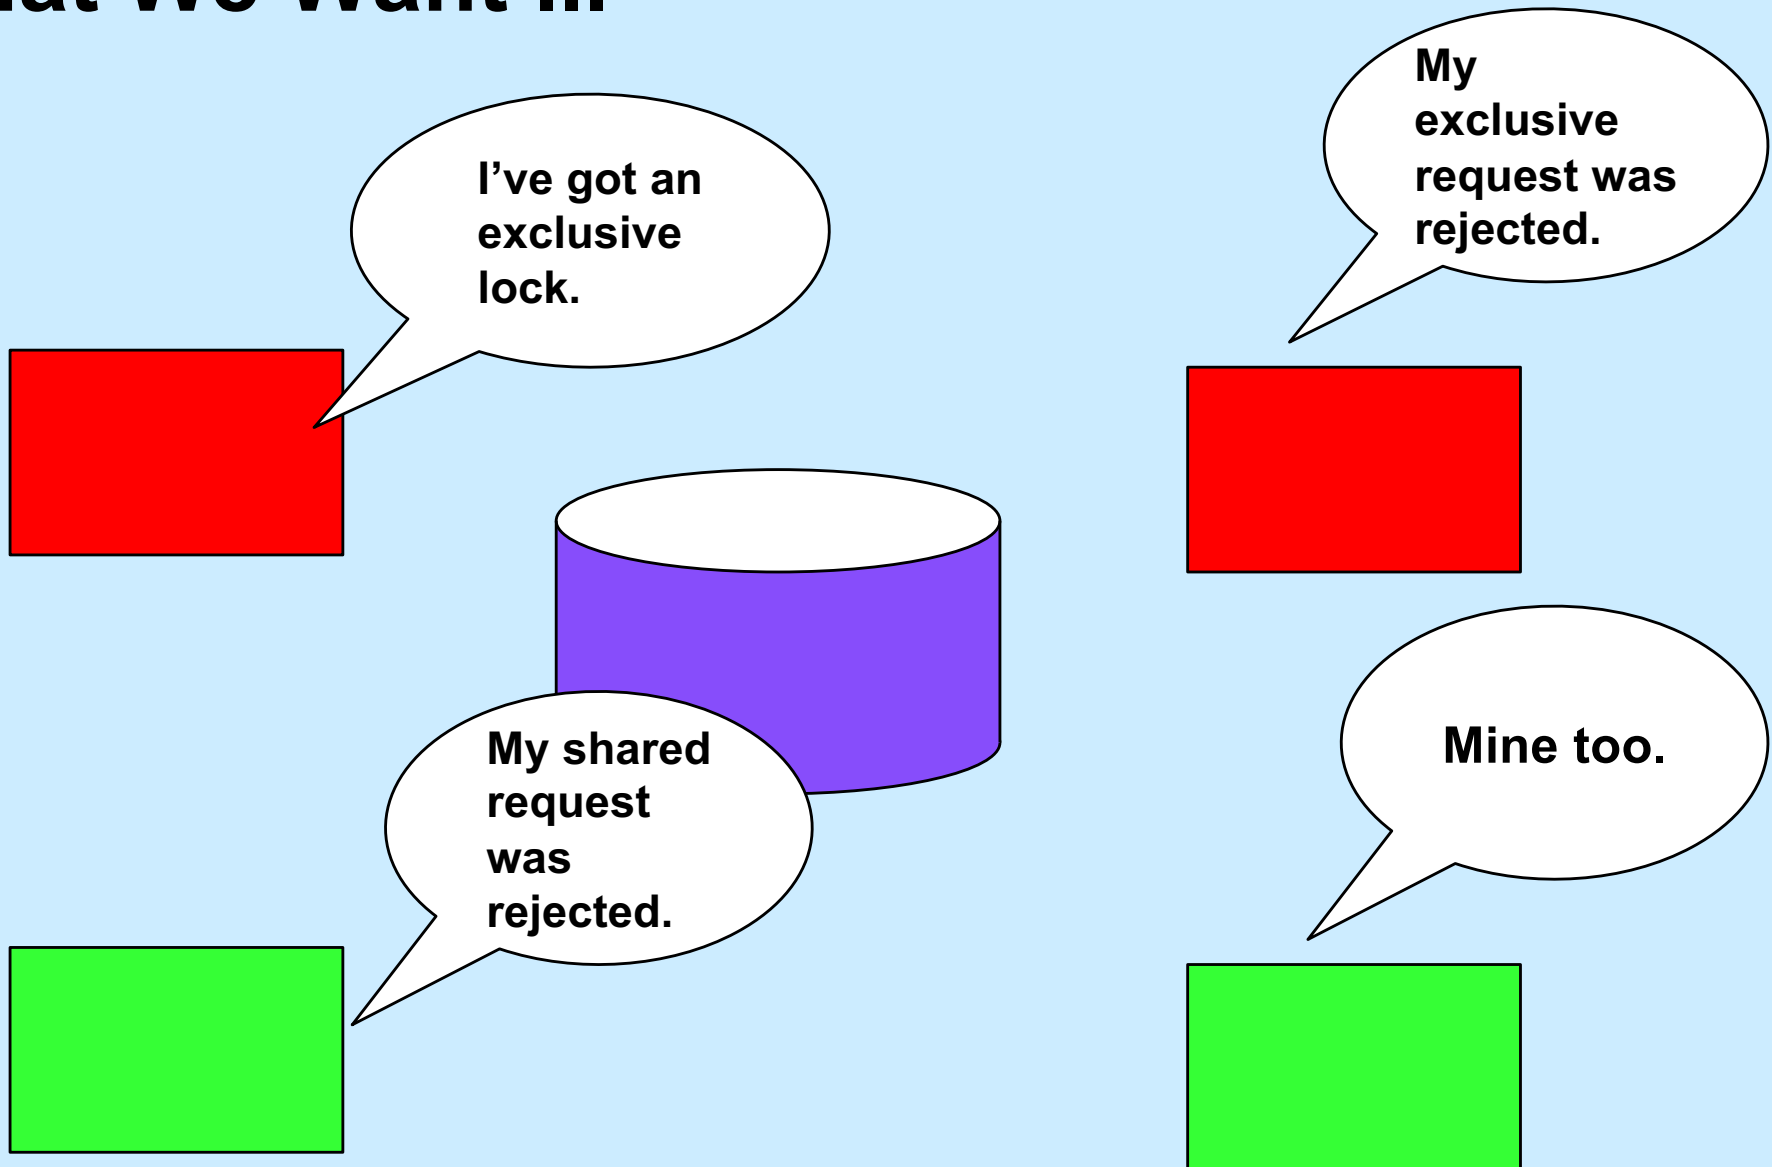
# Types of Locks

- **Shared (readers) locks**
  - any number may have them at same time
  - may not be held when an exclusive lock is held by others
- **Exclusive (writers) locks**
  - only one at a time
  - may not be held when any other lock is held by others

# What We Want ...



# What We Want ...



# Locking Files

- Early Unix didn't support file locking
- How did people survive?

- `open("file.lck", O_RDWR|O_CREAT|O_EXCL, 0666);`
    - » operation fails if *file.lck* exists, succeeds (and creates *file.lck*) otherwise
    - » requires cooperative programs

# Locking Files (continued)

- **How it's done in “modern” Unix**
  - “advisory locks” may be placed on files
    - » may request shared (readers) or exclusive (writers) lock
      - *fcntl* system call
    - » either succeeds or fails
    - » *open*, *read*, *write* always work, regardless of locks
    - » a lock applies to a specified range of bytes, not necessarily to the whole file
    - » requires cooperative programs
  - “mandatory locks” supported as a per-file option
    - » set along with permission bits
    - » if set, file can't be used unless process possesses appropriate locks



# Locking Files (still continued)

- **How to:**

```
struct flock fl;
fl.l_type = F_RDLCK;           // read lock
// fl.l_type = F_WRLCK;       // write lock
// fl.l_type = F_UNLCK;       // unlock
fl.l_whence = SEEK_SET;        // starting where
fl.l_start = 0;                // offset
fl.l_len = 0;                  // how much? (0 = whole file)
fd = open("file", O_RDWR);
if (fcntl(fd, F_SETLK, &fl) == -1)
    if ((errno == EACCES) || (errno == EAGAIN))
        // didn't get lock
    else
        // something else is wrong
else
    // got the lock!
```

---

# Quiz 2

- Your program currently has a shared lock on a portion of a file. It would like to “upgrade” the lock to be an exclusive lock. Would there be any problems with adding an option to *fcntl* that would allow the holder of a shared lock to wait until it’s possible to upgrade to an exclusive lock, then do the upgrade?
  - a) at least one major problem
  - b) either no problems whatsoever or some easy-to-deal-with problems

# Shell 1: Artisanal Coding

```
while ((line = get_a_line()) != 0) {
    tokens = parse_line(line);
    for (int i=0; i < ntokens; i++) {
        if (strcmp(tokens[i], ">") == 0) {
            // handle output redirection
        } else if (strcmp(tokens[i], "<") == 0) {
            // handle input redirection
        } else if (strcmp(tokens[i], "&") == 0) {
            // handle "no wait"
        } ... else {
            // handle other cases
        }
    }
    if (fork() == 0) {
        // ...
        execv(...);
    }
    // ...
}
```

# Shell 1: Non-Artisanal Coding (1)

```
while ((line = get_a_line()) != 0) {  
    tokens = parse_line(line);  
    for (int i=0; i < ntokens; i++) {  
        // handle "normal" case  
    }  
    if (fork() == 0) {  
        // ...  
        execv(...);  
    }  
    // ...  
}
```

# Shell 1: Non-Artisanal Coding (2)

```
next_line: while ((line = get_a_line()) != 0) {
    tokens = parse_line(line);
    for (int i=0; i < ntokens; i++) {
        if (redirection_symbol(token[i])) {
            // ...
            if (fork() == 0) {
                // ...
                execv(...); whoops!
            }
            // ...
            goto next_line;
        }
        // handle "normal" case
    }
    if (fork() == 0) {
        // ... (whoops!)
        execv(...);
    }
    // ...
}
```

# Shell 1: Non-Artisanal Coding (3)

```
next_line: while ((line = get_a_line()) != 0) {
    tokens = parse_line(line);
    for (int i=0; i < ntokens; i++) {
        if (redirection_symbol(token[i])) {
            // ...
            if (fork() == 0) {
                // ...
                execv(...);
            }
            // ... deal with &
            goto next_line;
        }
        // handle "normal" case
    }
    if (fork() == 0) {
        // ...
        execv(...);
    }
    // ... also deal with & here!
}
```

---

# Shell 1: Non-Artisanal Coding (Worse)

```
next_line: while ((line = get_a_line()) != 0) {  
    tokens = parse_line(line);  
    for (int i=0; i < ntokens; i++) {  
        if (redirection_symbol(token[i])) {  
            // ...  
            if (fork() == 0) {  
                // ...  
                execv(...);  
            }  
            // ... deal with &  
            goto next_line;  
        }  
        // handle "normal" case  
    }  
    if (fork() == 0) {  
        // ...  
        execv(...);  
    }  
    // ... also deal with & here!  
}
```

---

# Artisanal Programming

- **Factor your code!**
  - `A; FE | B; FE | C; FE = (A | B | C); FE`
- **Format as you write!**
  - don't run the formatter only just before handing it in
  - your code should always be well formatted
- **If you have a tough time understanding your code, you'll have a tougher time debugging it and TAs will have an even tougher time helping you**



# It's Your Code

- **Be proud of it!**
  - it not only works; it shows skillful artisanship
- **It's not enough to merely work**
  - others have to understand it
    - » (not to mention you ...)
  - you (and others) have to maintain it
    - » shell 2 is coming soon!