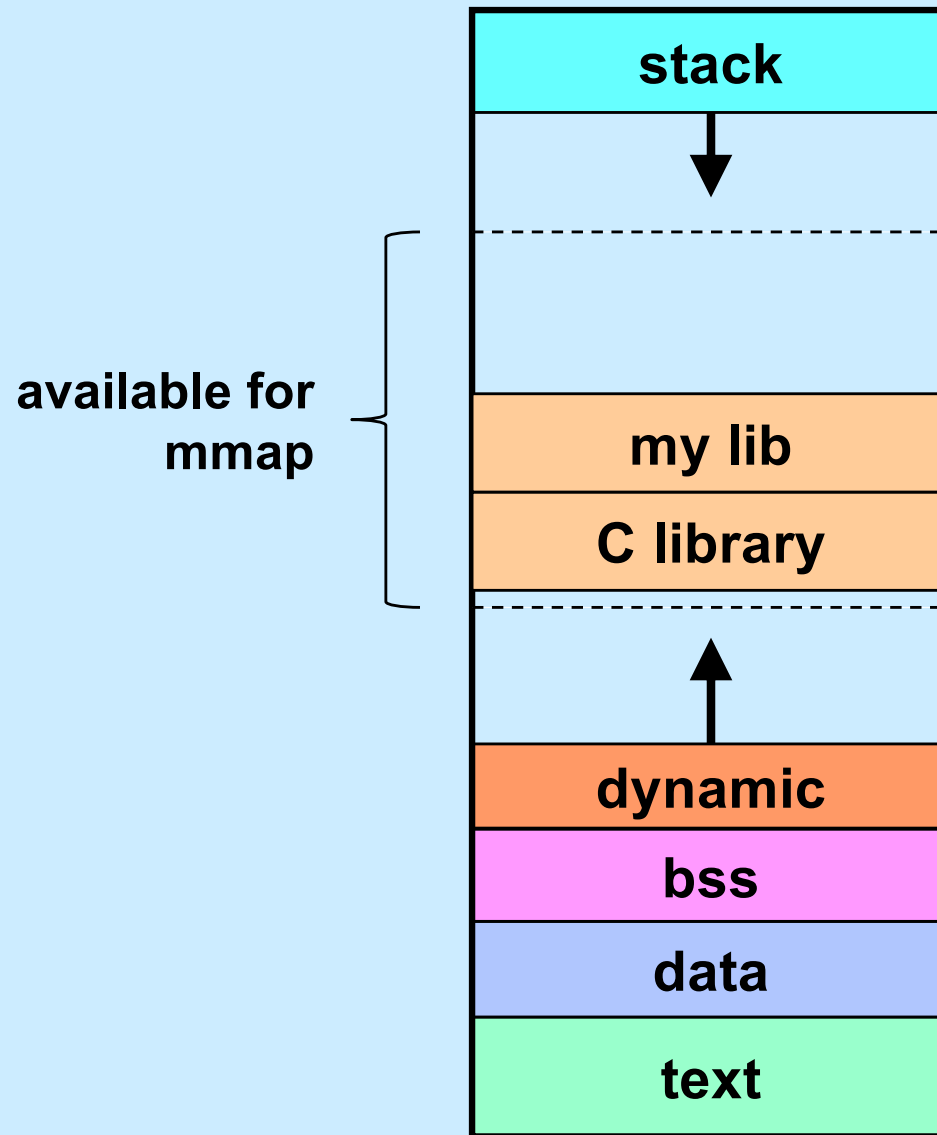


CS 33

Linking and Libraries (2)


Mmapping Libraries



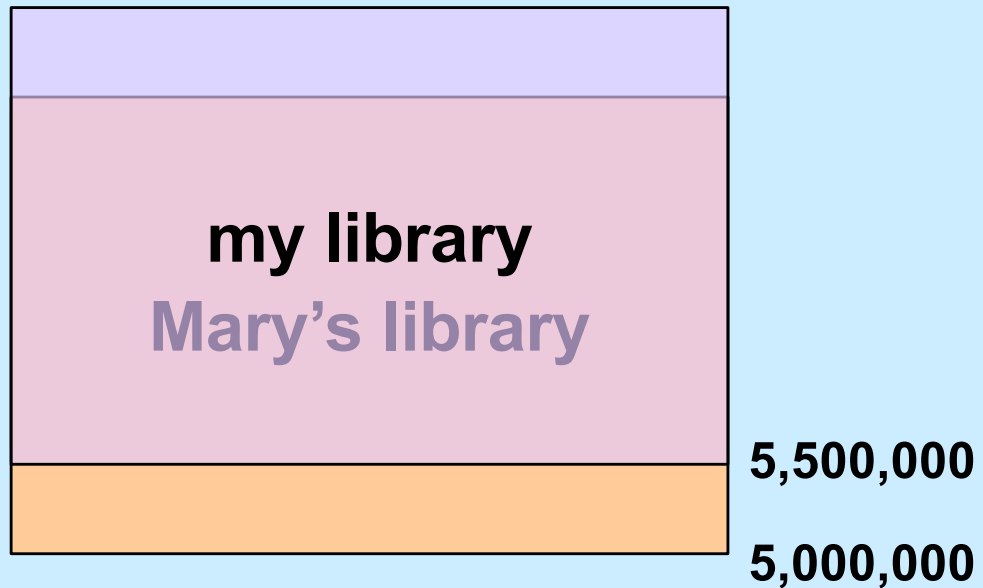
Problem

- **How is relocation handled?**

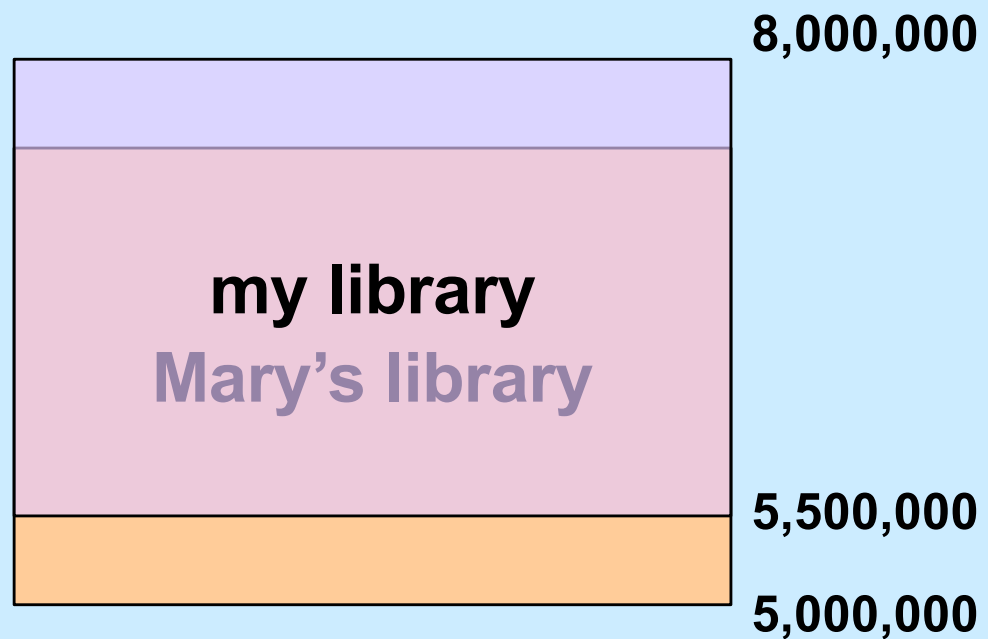
Pre-Relocation



But ...



But ...



Quiz 1

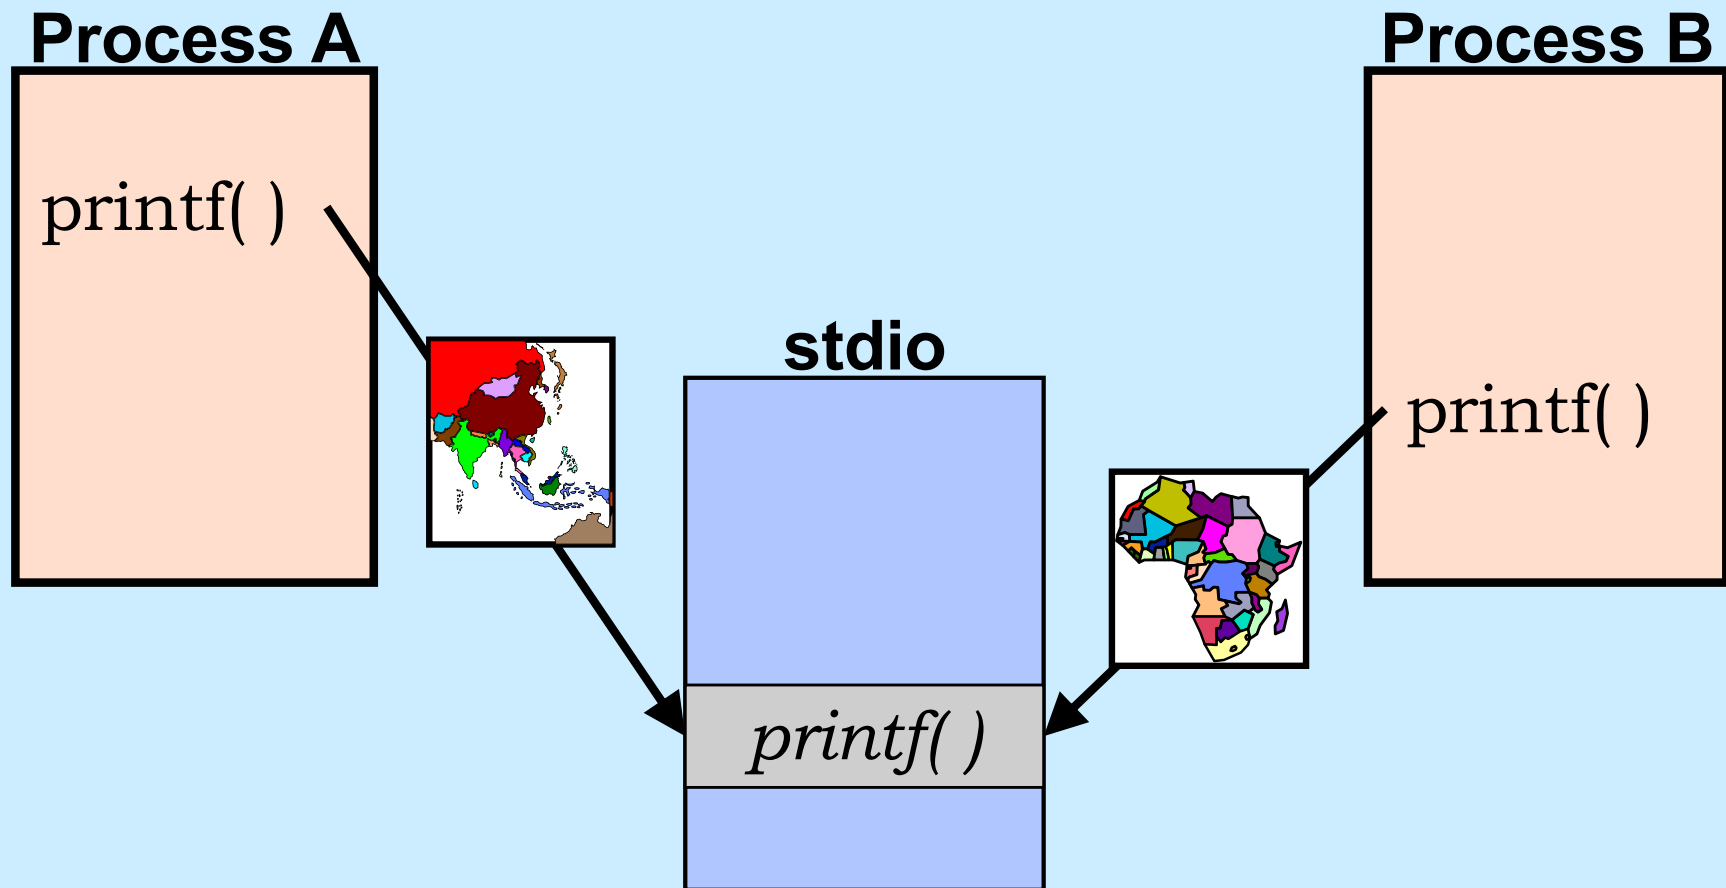
We need to relocate all references to Mary's library in my library. What option should we give to *mmap* when we map my library into our address space?

- a) the MAP_PRIVATE option**
- b) the MAP_SHARED option**
- c) mmap can't be used in this situation**

Relocation Revisited

- **Modify shared code to effect relocation**
 - result is no longer shared!
- **Separate shared code from (unshared) addresses**
 - position-independent code (PIC)
 - code can be placed anywhere
 - addresses in separate private section
 - » pointed to by a register

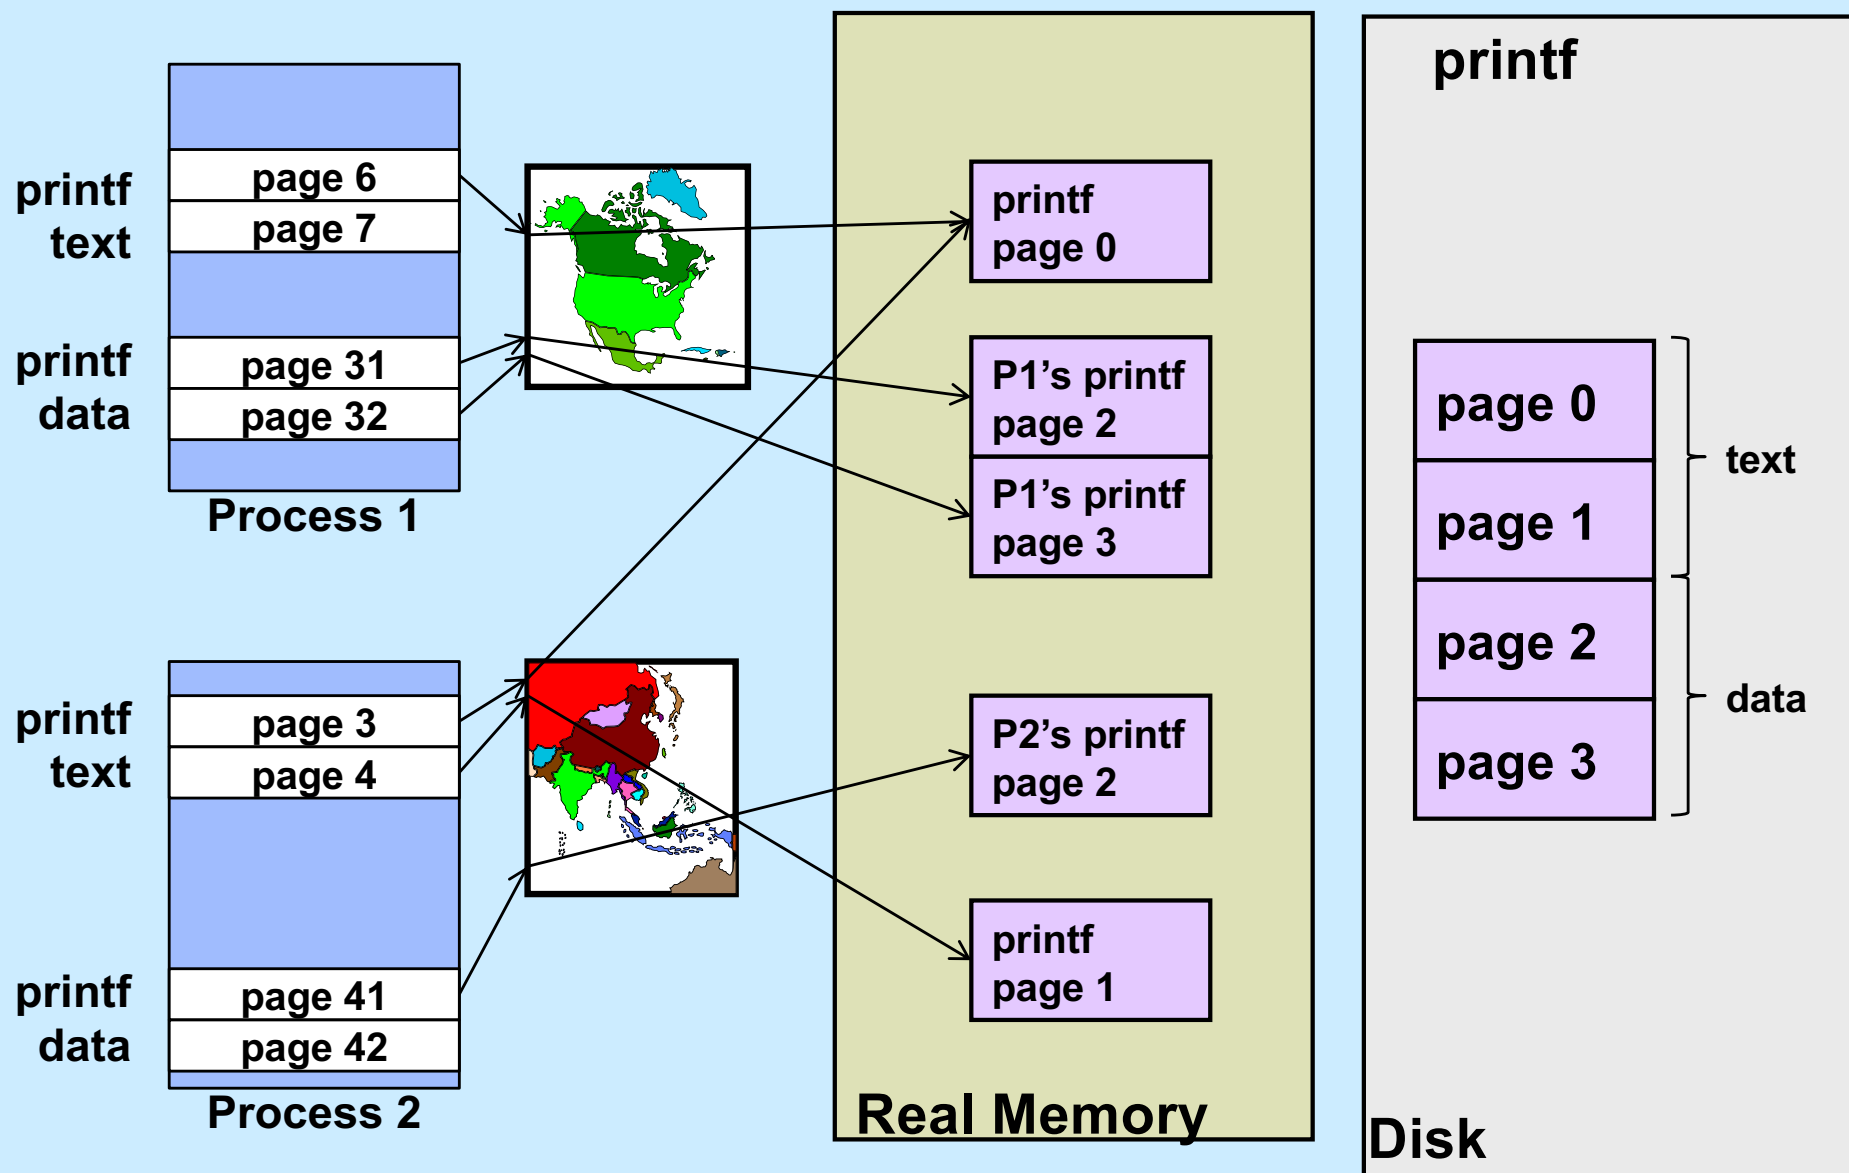
Mapping Shared Objects



Mapping printf into the Address Space

- **Printf's text**
 - read-only
 - can it be shared?
 - » yes: use MAP_SHARED
- **Printf's data**
 - read-write
 - not shared with other processes
 - initial values come from file
 - can mmap be used?
 - » MAP_SHARED wouldn't work
 - changes made to data by one process would be seen by others
 - » MAP_PRIVATE does work!
 - mapped region is initialized from file
 - changes are private

Mapping printf




Position-Independent Code

- **Produced by gcc when given the `-fPIC` flag**
- **Processor-dependent; x86-64:**
 - **each dynamic executable and shared object has:**
 - » **procedure-linkage table**
 - **shared, read-only executable code**
 - **essentially stubs for calling functions**
 - » **global-offset table**
 - **private, read-write data**
 - **relocated dynamically for each process**
 - » **relocation table**
 - **shared, read-only data**
 - **contains relocation info and symbol table**

Global-Offset Table: Data References

Global Offset Table →



Functions in Shared Objects

- Lots of them
- Many are never used
- Fix up linkages on demand

An Example

```
int main( ) {  
    puts("Hello world\n");  
    ...  
    return 0;  
}
```

00000000000000006b0 <main>:

6b0:	55	push	%rbp
6b1:	48 89 e5	mov	%rsp,%rbp
6b4:	48 8d 3d 99 00 00 00	lea	0x99(%rip),%rdi
6bb:	e8 a0 fe ff ff	callq	560 <puts@plt>

...

Before Calling puts

```
.PLT0:  
    pushq GOT+8(%rip)  
    jmp    *GOT+16(%rip)  
    nop; nop  
    nop; nop  
.puts:  
    jmp    *puts@GOT(%rip)  
.putsnext  
    pushq $putsRelOffset  
    jmp    .PLT0  
.PLT2:  
    jmp    *name2@GOT(%rip)  
.PLT2next  
    pushq $name2RelOffset  
    jmp    .PLT0
```

Procedure-Linkage Table

```
GOT:  
    .quad _DYNAMIC  
    .quad identification  
    .quad ld-linux.so  
  
puts:  
    .quad .putsnext  
name2:  
    .quad .PLT2next
```

Relocation info:

GOT_offset(puts) , symx(puts)

GOT_offset(name2) , symx(name2)

Relocation Table

After Calling puts

```
.PLT0:  
    pushq GOT+8(%rip)  
    jmp    *GOT+16(%rip)  
    nop; nop  
    nop; nop  
.puts:  
    jmp    *puts@GOT(%rip)  
.putsnext  
    pushq $putsRelOffset  
    jmp    .PLT0  
.PLT2:  
    jmp    *name2@GOT(%rip)  
.PLT2next  
    pushq $name2RelOffset  
    jmp    .PLT0
```

Procedure-Linkage Table

```
GOT:  
    .quad _DYNAMIC  
    .quad identification  
    .quad ld-linux.so  
  
puts:  
    .quad puts  
name2:  
    .quad .PLT2next
```

Relocation info:

GOT_offset(puts) , symx(puts)

GOT_offset(name2) , symx(name2)

Relocation Table

Not a Quiz!

On the second and subsequent calls to *puts*

- a) control goes directly to *puts*
- b) control goes to an instruction that jumps to *puts*
- c) control still goes to *ld-linux.so*, but it now transfers control directly to *puts*


CS 33

Multithreaded Programming (1)

Multithreaded Programming


- **A thread is a virtual processor**
 - an independent agent executing instructions
- **Multiple threads**
 - multiple independent agents executing instructions

Why Threads?



- Many things are easier to do with threads
- Many things run faster with threads

A Simple Example



Life Without Threads

```
void relay(int left, int right) {
    fd_set rd, wr;
    int left_read = 1, right_write = 0;
    int right_read = 1, left_write = 0;
    int sizeLR, sizeRL, wret;
    char bufLR[BSIZE], bufRL[BSIZE];
    char *bufpR, *bufpL;
    int maxFD = max(left, right) + 1;

    fcntl(left, F_SETFL, O_NONBLOCK);
    fcntl(right, F_SETFL, O_NONBLOCK);

    while(1) {
        FD_ZERO(&rd);
        FD_ZERO(&wr);
        if (left_read)
            FD_SET(left, &rd);
        if (right_read)
            FD_SET(right, &rd);
        if (left_write)
            FD_SET(left, &wr);
        if (right_write)
            FD_SET(right, &wr);


        select(maxFD, &rd, &wr, 0, 0);
```

```
        if (FD_ISSET(left, &rd)) {
            sizeLR = read(left, bufLR, BSIZE);
            left_read = 0;
            right_write = 1;
            bufpR = bufLR;
        }
        if (FD_ISSET(right, &rd)) {
            sizeRL = read(right, bufRL, BSIZE);
            right_read = 0;
            left_write = 1;
            bufpL = bufRL;
        }
        if (FD_ISSET(right, &wr)) {
            if ((wret = write(right, bufpR, sizeLR)) == sizeLR) {
                left_read = 1; right_write = 0;
            } else {
                sizeLR -= wret; bufpR += wret;
            }
        }
        if (FD_ISSET(left, &wr)) {
            if ((wret = write(left, bufpL, sizeRL)) == sizeRL) {
                right_read = 1; left_write = 0;
            } else {
                sizeRL -= wret; bufpL += wret;
            }
        }
    }
    return 0;
}
```


Life With Threads

```
void copy(int source, int destination) {  
    struct args *targs = args;  
    char buf[BSIZE];  
  
    while(1) {  
        int len = read(source, buf, BSIZE);  
        write(destination, buf, len);  
    }  
}
```



Processes vs. Threads



Process 1





Process 2



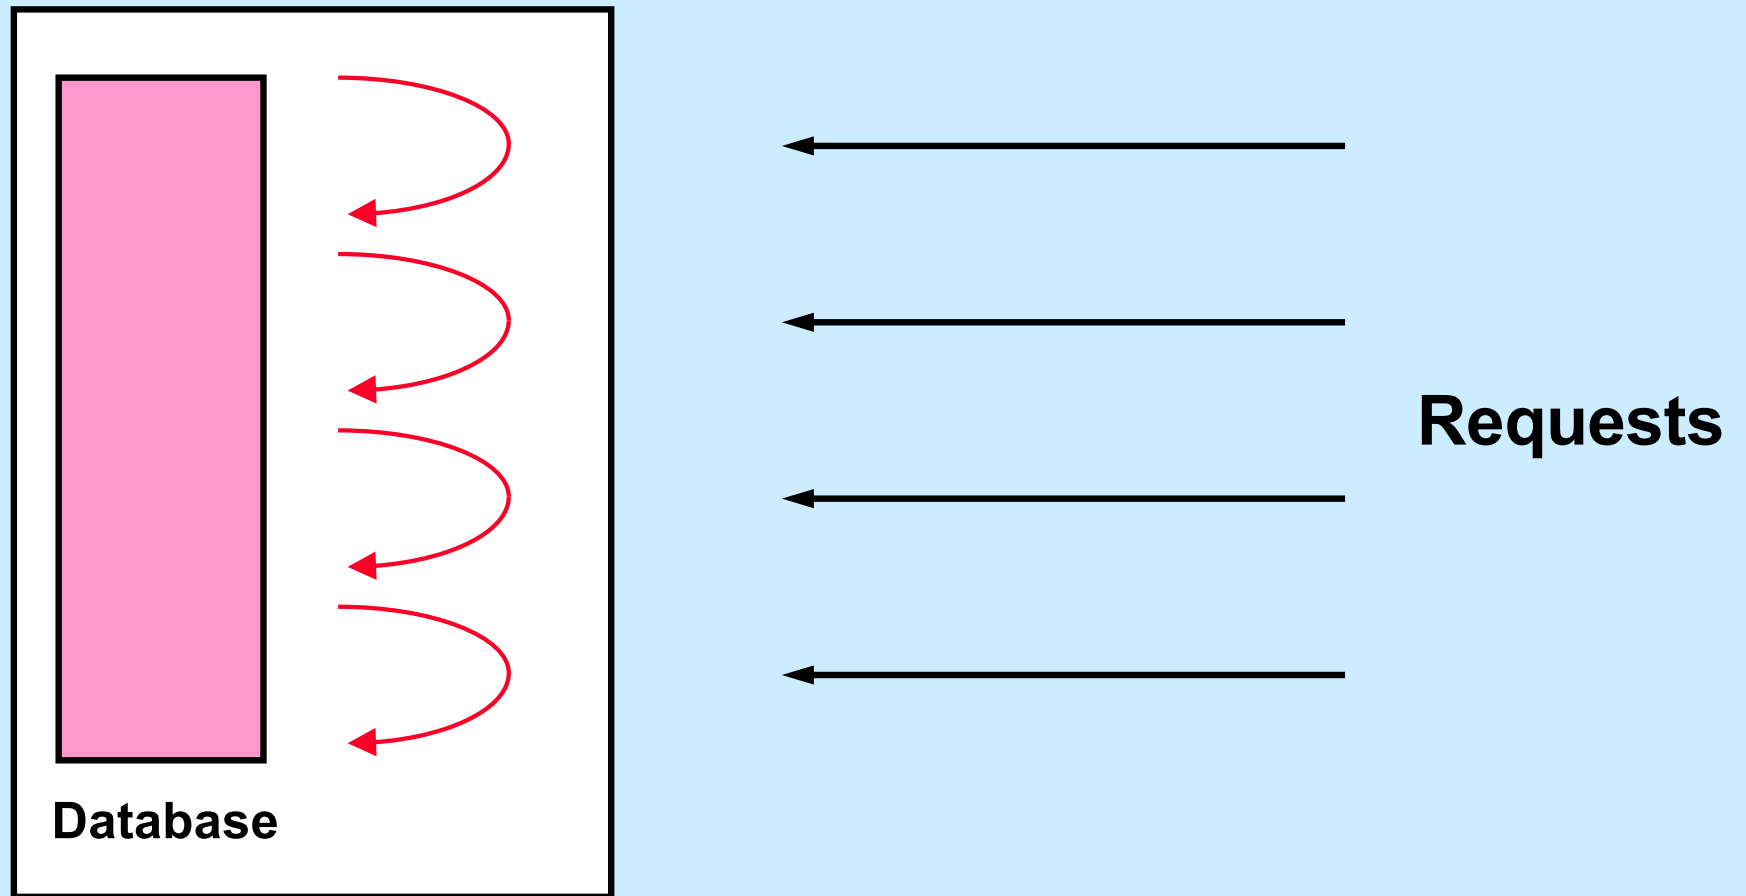
Process 3

Single-Threaded Database Server




Requests


Multithreaded Database Server




Single-Core Chips



Dual-Core Chips



Multi-Core Chips



Good News/Bad News


Good news

- multi-threaded programs can take advantage of multi-core chips (single-threaded programs cannot)

Bad news

- it's not easy
 - » must have parallel algorithm
 - employing at least as many threads as processors
 - threads must keep processors busy
 - doing useful work

Matrix Multiplication Revisited



Standards

- **POSIX 1003.4a → 1003.1c → 1003.1j**
- **Microsoft**
 - **Win32/64**

Creating Threads

```
long A[M][N], B[N][P], C[M][P];  
...  
for (i=0; i<M; i++)    // create worker threads  
    pthread_create(&thr[i], 0, matmult, i);  
  
...
```

```
void *matmult(void *arg) {  
    long i = (long)arg;  
    // compute row i of the product C of A and B  
    ...  
}
```

When Is It Finished?

```
long A[M][N], B[N][P], C[M][P];  
...  
for (i=0; i<M; i++)    // create worker threads  
    pthread_create(&thr[i], 0, matmult, i));  
  
for (i=0; i<M; i++)    // wait for termination  
    pthread_join(thr[i], 0);  
  
printResult(C); // shouldn't do this until  
                // workers have terminated
```

Example (1)

```
#include <stdio.h>
#include <pthread.h>
#include <string.h>
```

```
#define M    3
#define N    4
#define P    5
```

```
long A[M][N];
long B[N][P];
long C[M][P];
```

```
void *matmult(void *);
```

```
main( ) {
    long i;
    pthread_t thr[M];
    int error;

    // initialize the matrices
    ...
}
```

Example (2)

```
for (i=0; i<M; i++) { // create worker threads
    if (error = pthread_create(
        &thr[i],
        0,
        matmult,
        (void *)i)) {
        fprintf(stderr, "pthread_create: %s", strerror(error));
        exit(1);
    }
}

for (i=0; i<M; i++) // wait for workers to finish their jobs
    pthread_join(thr[i], 0)

/* print the results ... */
}
```

Example (3)

```
void *matmult(void *arg) {  
    long row = (long) arg;  
    long col;  
    long i;  
    long t;  
  
    for (col=0; col < P; col++) {  
        t = 0;  
        for (i=0; i<N; i++)  
            t += A[row][i] * B[i][col];  
        C[row][col] = t;  
    }  
    return (0);  
}
```

Compiling It

```
% gcc -o mat mat.c -pthread
```

Termination

```
pthread_exit((void *) value);
```

```
return((void *) value);
```

```
pthread_join(thread, (void **) &value);
```


Detached Threads

```
start_servers( ) {  
    pthread_t thread;  
    int i;  
  
    for (i=0; i<nr_of_server_threads; i++) {  
        pthread_create(&thread, 0, server, 0);  
        pthread_detach(thread);  
    }  
    ...  
}  
  
void *server(void * arg) {  
    ...  
}
```

Complications

```
void relay(int left, int right) {  
    pthread_t LRthread, RLthread;  
  
    pthread_create(&LRthread,  
        0,  
        copy,  
        left, right);           // Can't do this ...  
    pthread_create(&RLthread,  
        0,  
        copy,  
        right, left);           // Can't do this ...  
}
```

Multiple Arguments

```
typedef struct args {  
    int src;  
    int dest;  
} args_t;
```

```
void relay(int left, int right) {  
    args_t LRargs, RLargs;  
    pthread_t LRthread, RLthread;  
    ...  
    pthread_create(&LRthread, 0, copy, &LRargs);  
    pthread_create(&RLthread, 0, copy, &RLargs);  
}
```

Multiple Arguments

```
typedef struct args {  
    int src;  
    int dest;  
} args_t;
```

```
void relay(int left, int right) {  
    args_t LRargs, RLargs;  
    pthread_t LRthread, RLthread;  
    ...  
    pthread_create(&LRthread, 0, copy, &LRargs);  
    pthread_create(&RLthread, 0, copy, &RLargs);  
}
```


Quiz 2

Does this work?


a) yes

b) no

Execution



Multiplexing Processors



Quiz 3

```
pthread_create(&tid, 0, tproc, (void *)1);  
pthread_create(&tid, 0, tproc, (void *)2);
```

```
printf("T0\n");
```

```
...
```

```
void *tproc(void *arg) {  
    printf("T%d\n", (long) arg);  
    return 0;  
}
```

In which order are things printed?

- a) T0, T1, T2
- b) T1, T2, T0
- c) T2, T1, T0
- d) indeterminate

Cost of Threads

```
int main(int argc, char *argv[]) {
    ...
    val = niters/nthreads;

    for (i=0; i<nthreads; i++)
        pthread_create(&thread, 0, work, (void *)val);
    pthread_exit(0);
    return 0;
}

void *work(void *arg) {
    long n = (long)arg; int i, j; volatile long x;

    for (i=0; i<n; i++) {
        x = 0;
        for (j=0; j<1000; j++)
            x = x*j;
    }
    return 0;
}
```


Cost of Threads

```
int main(int argc, char *argv[]) {
    ...
    val = niters/nthreads;

    for (i=0; i<nthreads; i++)
        pthread_create(&thread, 0, work, (void *)val);
    pthread_exit(0);
    return 0;
}

void *work(void *arg) {
    long n = (long)arg; int i, j; volatile long x;

    for (i=0; i<n; i++) {
        x = 0;
        for (j=0; j<1000; j++)
            x = x*j;
    }
    return 0;
}
```

Quiz 4

This code runs in time n on a 4-core processor when $nthreads$ is 8. It runs in time p on the same processor when $nthreads$ is 400.

- a) $n \ll p$ (slower)
- b) $n \approx p$ (same speed)
- c) $n \gg p$ (faster)

Problem

```
pthread_create(&thread, 0, start, 0);
```

```
...
```

```
void *start(void *arg) {  
    long BigArray[128*1024*1024];  
    ...  
    return 0;  
}
```

Thread Attributes

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
  
...  
  
/* establish some attributes */  
  
...  
  
pthread_create(&thread, &thr_attr, startroutine, arg);  
  
...
```

Stack Size

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
pthread_attr_setstacksize(&thr_attr, 130*1024*1024);  
  
...  
  
pthread_create(&thread, &thr_attr, startroutine, arg);  
  
...
```